

Class Activity 3 [25 minutes]: normal-gamma toy model

Last we left out intrepid heros, we desired samples from

The Full (Joint) Posterior Distribution	Full Conditional Distributions
$p(\theta, \phi x) = p(\theta, \phi x, \theta_0, \tau, \alpha, \beta)$	$p(\theta x, \theta_0, \tau, \phi) = N\left(\frac{(\tau\theta_0 + \phi \sum_{i=1}^n x_i)}{(\tau + n\phi)}, \sigma^{-2} = \tau + n\phi\right)$ $p(\phi x, \alpha, \beta, \theta) = \text{Gamma}\left(\alpha + \frac{n}{2}, \beta + \frac{1}{2} \sum_{i=1}^n (x_i - \theta)^2\right)$
But had only derived	

Alternatingly sample, back and fourth, between the two full conditionals, conditioning the sample draw for the current parameter on the last drawn value of the other parameter...

0. You'll need data... what's the data?

1. You'll need **actual parameters values** to simulate data... what are θ_{true} and ϕ_{true} ?

2. You'll need to know your **prior hyperparameters**... what are θ_0 , τ , α , and β ?

3. You'll need **initial values** from which to start the process... what are $\theta^{(0)}$ or $\phi^{(0)}$?

4. You'll need to sample from the **full conditional distributions**... how to do that?

5. You'll need to visualize the $\theta^{(i)}$ and $\phi^{(i)}$ samples for $i = 1, \dots, T$... how to do that?

6. You'll need to summarise the $\theta^{(i)}$ and $\phi^{(i)}$ samples for $i = 1, \dots, T$... how to do that?

Class Activity 3: Solution [5 minutes]

```
[2]: import numpy as np; from scipy import stats; import matplotlib.pyplot as plt


theta_0,tau = 0,1; alpha,beta = 2,1/2 # Initialize Prior Hyperparameters:
    ↳ specify prior distributions
theta_prior = stats.norm(loc=theta_0, scale=tau) # https://en.wikipedia.org/wiki/
    ↳ Gamma_distribution
phi_prior = stats.gamma(a=alpha, scale=1/beta) # Note: scale is 1/ beta (beta
    ↳ is rate parameter)

# In real world we probably wouldn't know for sure what distribution our data x
    ↳ is from:
# We wouldn't know the theta_true / phi_true, or even if the likelihood was
    ↳ normally distributed...
# and don't forget about the priors/hyperparameters above... these also require
    ↳ considerations...
theta_true,phi_true = 1,1; true_data_model = stats.
    ↳ norm(loc=theta_true,scale=phi_true)
n = 50; np.random.seed(seed=123); x = true_data_model.rvs(n);
```

```

Rplus = np.linspace(0,15,100); Rpm = np.linspace(-3,5,100)
fig,ax = plt.subplots(1,3,figsize=(18,5)); ax[0].plot(Rpm, theta_prior.pdf(Rpm));
→ ax[0].set_title('$\\theta$ prior'); ax[0].vlines(theta_true,0,.
→45,'k',label='True $\\theta$'); ax[0].legend(); ax[1].plot(Rplus, phi_prior.
→pdf(Rplus)); ax[1].vlines(phi_true,0,.2,'k',label='True $\\phi$'); ax[1].
→legend(); ax[1].set_title('$\\phi$ prior'); ax[2].hist(x, density=True); ax[2].
→plot(Rpm, true_data_model.pdf(Rpm)); ax[2].set_title('Observed Data and True
→Data Distribution');

```



Class Activity 3: Solution [5 minutes]

```

[5]: # Solution from Leo Watson
iterations = 100 # Gibbs Samples
theta_gibbs_posterior_mcmc_draws, phi_gibbs_posterior_mcmc_draws = np.
→ zeros(iterations+1), np.zeros(iterations+1)
# initial_theta, initial_phi = 0, 0 # Not really an initial phi as we'll do the
→ phi first per update
# Ideally we want our initial values near the "typical set" which along with
→ similar density values
# comprise the large majority of the probability of the distribution so the
→ convergence is faster
np.random.seed(123)
for i in range(1, iterations + 1):
    theta_old = theta_gibbs_posterior_mcmc_draws[i-1]

    phi_new = stats.gamma(a = alpha+n/2, scale = (1/(beta+1/2*((x-theta_old)**2).
→ sum()))).rvs()
    phi_gibbs_posterior_mcmc_draws[i] = phi_new


    theta_new = stats.norm(loc= (tau*theta_0+phi_new*sum(x))/(tau+n*phi_new),
→ scale= 1/np.sqrt(tau+n*phi_new)).rvs()
    theta_gibbs_posterior_mcmc_draws[i] = theta_new

```

```

burn = 20
fig,ax = plt.subplots(1,4, figsize=(18,4)); ax[0].
    plot(theta_gibbs_posterior_mcmc_draws[:burn],label="$\\theta$ burn in"); ax[1].
    plot(theta_gibbs_posterior_mcmc_draws[burn:],label="$\\theta$ MCMC Posterior_
    Draws"); ax[2].plot(phi_gibbs_posterior_mcmc_draws[:burn], label="$\\phi$ burn_
    in"); ax[3].plot(phi_gibbs_posterior_mcmc_draws[burn:], label="$\\phi$ MCMC_
    Posterior Draws"); ax[1].hlines(theta_true,0,iterations,'k',label='True_
    $\\theta$'); ax[3].hlines(phi_true,0,iterations,'k',label='True $\\phi$');
    ax[1].hlines(x.mean(),0,iterations,'r',linestyle='dashed',label='Sample Mean');
    ax[1].legend(); ax[3].hlines(1/x.
    var(),0,iterations,'r',linestyle='dashed',label='Sample Precision'); ax[3].
    legend(); ax[0].legend(); ax[2].legend();

```



Class Activity 3: Solution [3 minutes]

```

[6]: # Solution from Leo Watson # DEPENDENCE GETS WASHED OUT AND PROBABLY STOPS_
    MATTERING IN THE LONG RUN
iterations = 1020 # Gibbs Samples
theta_gibbs_posterior_mcmc_draws,phi_gibbs_posterior_mcmc_draws = np.
    zeros(iterations+1),np.zeros(iterations+1)
# initial_theta,initial_phi = 0,0 # Not really an initial phi as we'll do the_
    phi first per update
# Ideally we want our initial values near the "typical set" which along with_
    similar density values
# comprise the large majority of the probability of the distribution so the_
    convergence is faster
np.random.seed(123)
for i in range(1, iterations + 1):
    theta_old = theta_gibbs_posterior_mcmc_draws[i-1]

    phi_new = stats.gamma(a = alpha+n/2, scale = (1/(beta+1/
    2*sum((x-theta_old)**2))))).rvs()
    phi_gibbs_posterior_mcmc_draws[i] = phi_new

```

```

theta_new = stats.norm(loc= (tau*theta_0+phi_new*sum(x))/(tau+n*phi_new),
                        scale= 1/np.sqrt(tau+n*phi_new)).rvs()
theta_gibbs_posterior_mcmc_draws[i] = theta_new


```

```
burn = 20
```

```

fig,ax = plt.subplots(1,4, figsize=(18,4)); ax[0].
    →plot(theta_gibbs_posterior_mcmc_draws[:burn],label="$\\theta$ burn in"); ax[1].
    →plot(theta_gibbs_posterior_mcmc_draws[burn:],label="$\\theta$ MCMC Posterior_
    →Draws"); ax[2].plot(phi_gibbs_posterior_mcmc_draws[:burn], label="$\\phi$ burn_
    →in"); ax[3].plot(phi_gibbs_posterior_mcmc_draws[burn:], label="$\\phi$ MCMC_
    →Posterior Draws"); ax[1].hlines(theta_true,0,iterations,'k',label='True_
    →$\\theta$'); ax[3].hlines(phi_true,0,iterations,'k',label='True $\\phi$');
    →ax[1].hlines(x.mean(),0,iterations,'r',linestyle='dashed',label='Sample Mean');
    → ax[1].legend(); ax[3].hlines(1/x.
    →var(),0,iterations,'r',linestyle='dashed',label='Sample Precision'); ax[3].
    →legend(); ax[0].legend(); ax[2].legend();

```



Class Activity 3: Solution [4 minutes]

```


[8]: import numpy as np; from scipy import stats; import matplotlib.pyplot as plt
theta_0,tau = 0,1; alpha,beta = 2,1/2 # Initialize Prior Hyperparameters:
    →specify prior distributions
theta_prior = stats.norm(loc=theta_0, scale=tau) # https://en.wikipedia.org/wiki/
    →Gamma_distribution
phi_prior = stats.gamma(a=alpha, scale=1/beta) # Note: scale is 1/ beta (beta_
    →is rate parameter)
# In real world we probably wouldn't know for sure what distribution our data x_
    →is from:
# We wouldn't know the theta_true / phi_true, or even if the likelihood was_
    →normally distributed...
# and don't forget about the priors/hyperparameters above... these also require_
    →considerations...
theta_true,phi_true = 1,1; true_data_model = stats.
    →norm(loc=theta_true,scale=phi_true)
n = 50; np.random.seed(seed=123); x = true_data_model.rvs(n);

```

```

Rplus = np.linspace(0,15,100); Rpm = np.linspace(-3,3,100)
fig,ax = plt.subplots(1,3,figsize=(18,5)); ax[0].plot(Rpm, theta_prior.pdf(Rpm),
→label='$\\theta$ prior'); ax[0].set_title('$\\theta$ prior to posterior_
→update'); ax[0].vlines(theta_true,0,2,'k',label='True $\\theta$'); ax[1].
→plot(Rplus, phi_prior.pdf(Rplus), label='$\\phi$ prior'); ax[1].
→vlines(phi_true,0,2.5,'k',label='True $\\phi$'); ax[1].set_title('$\\phi$
→prior to posterior update'); ax[2].hist(x, density=True); ax[2].plot(Rpm,
→true_data_model.pdf(Rpm)); ax[2].set_title('Observed Data and True Data_
→Distribution');
from scipy.stats import gaussian_kde
theta_gibbs_posterior_mcmc_draws_sorted =
→sorted(theta_gibbs_posterior_mcmc_draws);
ax[0].
→plot(theta_gibbs_posterior_mcmc_draws_sorted,gaussian_kde(theta_gibbs_posterior_mcmc_draws_so
→label='$\\theta$ posterior'); ax[0].vlines(x.
→mean(),0,2,'r',linestyle='dashed',label='Sample Mean');
phi_gibbs_posterior_mcmc_draws_sorted = sorted(phi_gibbs_posterior_mcmc_draws);
ax[1].
→plot(phi_gibbs_posterior_mcmc_draws_sorted,gaussian_kde(phi_gibbs_posterior_mcmc_draws_sorted
→label='$\\phi$ posterior'); ax[1].vlines(1/x.var(),0,2.
→5,'r',linestyle='dashed',label='Sample Precision'); ax[0].legend(); ax[1].
→legend(); ax[0].set_xlim([-0.5,1.5]); ax[1].set_xlim([0,1.5]);

```



Gibbs Sampling: History [8 minutes]

- There is a [log-linear model](#) known as the [Boltzmann distribution](#).... > The *Boltzmann distribution* is also called the *Gibbs distribution* after the statistical physicist whose research frequently utilized and explored it
- In 1984 the *Geman brothers* figured out an algorithm to sample from the Boltzmann distribution that had analogy to a statistical physics processes analyzed by Gibbs, *so they named their sampler* after *Gibbs*...
- In 1990 [my first advisor from 2005-2007] Alan Gelfand introduced a generalized *Gibbs sampling* algorithm to the statistical community as a powerful tool for generalized *Bayesian*

posterior analysis...

Gibbs sampling initially dominated Bayesian posterior analysis; but, the cyclical “one dimension at a time” nature of **Gibbs sampling** suffers from a prohibitive “curse of dimensionality” limiting its application to low dimensional parameter spaces. **Hamiltonian Monte Carlo (HMC)** generally avoids this for **continuous** distributions, so **Gibbs sampling** is now only used in **discrete contexts** (often in conjunction with **Metropolis-Hastings**) where **HMC** can’t be used.

Gibbs Sampling: Algorithm [10 minutes]

The Gibbs Sampler is a simple algorithm...cyclically sampling through **Full Conditional Distributions** produces samples from the corresponding **Full (Joint) Posterior Distribution** just as you’ve done above for the class activity...

Gibbs sampling: cycle through **full conditionals**

1. $p(\theta_1^{(i)} | \theta_2^{(i-1)}, \dots, \theta_q^{(i-1)}, x)$
 2. $p(\theta_2^{(i)} | \theta_1^{(i-1)}, \dots, \theta_q^{(i-1)}, \theta_1^{(i)}, x)$
 3. $p(\theta_3^{(i)} | \theta_1^{(i-1)}, \dots, \theta_q^{(i-1)}, \theta_2^{(i)}, \theta_1^{(i)}, x)$
 4. $p(\theta_4^{(i)} | \theta_1^{(i-1)}, \dots, \theta_q^{(i-1)}, \theta_3^{(i)}, \theta_2^{(i)}, \theta_1^{(i)}, x)$
 - ...
 - q. $p(\theta_q^{(i)} | \theta_1^{(i-1)}, \dots, \theta_q^{(i-1)}, \theta_1^{(i)}, \dots, \theta_{q-1}^{(i)}, x)$
- completes one full cycle; return to 1...

This is NOT the chain rule...

$$\cancel{p(\theta_1^{(i)} | x) p(\theta_2^{(i)} | \theta_1^{(i)}, x) \dots p(\theta_q^{(i)} | \theta_{q-1}^{(i)}, \dots, \theta_2^{(i)}, \theta_1^{(i)}, x)}$$

This is a **Markov process** for (vector) θ since $\theta^{(i)}$ depends only on $\theta^{(i-1)}$.

The **Gibbs sampling Markov process** will converge to a **stationary distribution** which will be the **posterior distribution** $p(\theta | x)$

The samples from **Gibbs sampling** are a **Markov Chain** of sequentially (**Markov**) dependent samples forming a **Markov process** whose **stationary distribution** will be the **posterior distribution** $p(\theta | x)$

Markov Chain Monte Carlo (MCMC): Idea [4 minutes]

Bayesian posterior analysis is based on Markov Chain Monte Carlo - **Markov Chain**: sequentially dependent samples (from a stationary/posterior distribution) - **Monte Carlo (Integration)**: since **expectation** \equiv **integration** so **integration** \equiv **estimation**

This is because Bayesian posterior inference takes two forms:

$$\bullet \quad \bar{x} = \frac{1}{T} \sum_{t=1}^N x_i \approx E[X] = \int x f(x) dx$$

Markov Chain Monte Carlo (MCMC): Idea [6 minutes]

Bayesian posterior analysis is based on Markov Chain Monte Carlo - **Markov Chain**: sequentially dependent samples (from a stationary/posterior distribution) - **Monte Carlo (Integration)**: since **expectation** \equiv **integration** so **integration** \equiv **estimation**

This is because Bayesian posterior inference takes two forms:

$$\bullet \quad E[\theta | x] = \int \theta p(\theta | x) d\theta \quad \text{is estimated with} \quad \bar{\theta} = \frac{1}{T} \sum_{t=1}^T \theta^{(t)} \quad \text{for} \quad \theta^{(t)} \sim p(\theta | x)$$

- Estimation is an “average” of possible parameter values “weighted” by their uncertainty which is encoded in our posterior belief about the parameters... “fairly” takes into account all possibilities
- Whereas “best choice” MLE or $\theta^{\text{MAP}} = \max_{\theta} p(\theta|x)$ is sensitive to sample overfitting... MLE or θ^{MAP} just the best for the sample at hand but doesn’t take into account uncertainty...

Markov Chain Monte Carlo (MCMC): Idea [2 minutes]

Bayesian posterior analysis is based on **Markov Chain Monte Carlo** - *Markov Chain*: sequentially dependent samples (from a stationary/posterior distribution) - *Monte Carlo (Integration)*: since **expectation** \equiv **integration** so **integration** \equiv **estimation**

This is because Bayesian posterior inference takes two forms:

- $\Pr(\theta \leq a) = \int_{-\infty}^a p(\theta|x)d\theta$ is estimated with $\hat{p}_{\theta \leq a} = \frac{1}{T} \sum_{t=1}^T 1_{[\infty, a]}(\theta^{(t)})$ for $\theta^{(t)} \sim p(\theta|x)$
 - Estimation is an “average” of possible parameter values “weighted” by their uncertainty which is encoded in our posterior belief about the parameters... “fairly” takes into account all possibilities
 - Whereas “best choice” MLE or $\theta^{\text{MAP}} = \max_{\theta} p(\theta|x)$ is sensitive to sample overfitting... it’s just the best for the sample at hand but doesn’t take into account uncertainty...

Markov Chain Monte Carlo (MCMC): Idea [6 minutes]

Bayesian posterior analysis is based on **Markov Chain Monte Carlo** - *Markov Chain*: sequentially dependent samples (from a stationary/posterior distribution) - *Monte Carlo (Integration)*: since **expectation** \equiv **integration** so **integration** \equiv **estimation**

This is because Bayesian posterior inference takes two forms:

- $(1 - \alpha) \approx \int_a^b p(\theta|x)d\theta$ for $a = \theta_{(\lfloor T \times \frac{\alpha}{2} \rfloor)}$ and $b = \theta_{(\lceil T \times (1 - \frac{\alpha}{2}) \rceil)}$ and $\theta^{(t)} \sim p(\theta|x)$
e.g., $\alpha = 0.05$ where $\theta_{(k)}$ is the k^{th} rank ordered posterior sample
 - Interval estimation is the range of possible parameter values we believe might be plausible based on our posterior uncertainty beliefs... given as the probability statement $\Pr_{\theta|x}(a \leq \theta \leq b) = 1 - \alpha$
 - This is a $100 \times (1 - \alpha)\%$ **Bayesian credible interval** conveying posterior belief about parameter θ
Bayesian posterior inference remains based upon **integration** rather than **maximization**

Markov Chain Monte Carlo (MCMC): Idea [2 minutes]

Bayesian posterior analysis is based on Markov Chain Monte Carlo - *Markov Chain*: sequentially dependent samples (from a stationary/posterior distribution) - *Monte Carlo (Integration)*: since **expectation** \equiv **integration** so **integration** \equiv **estimation**

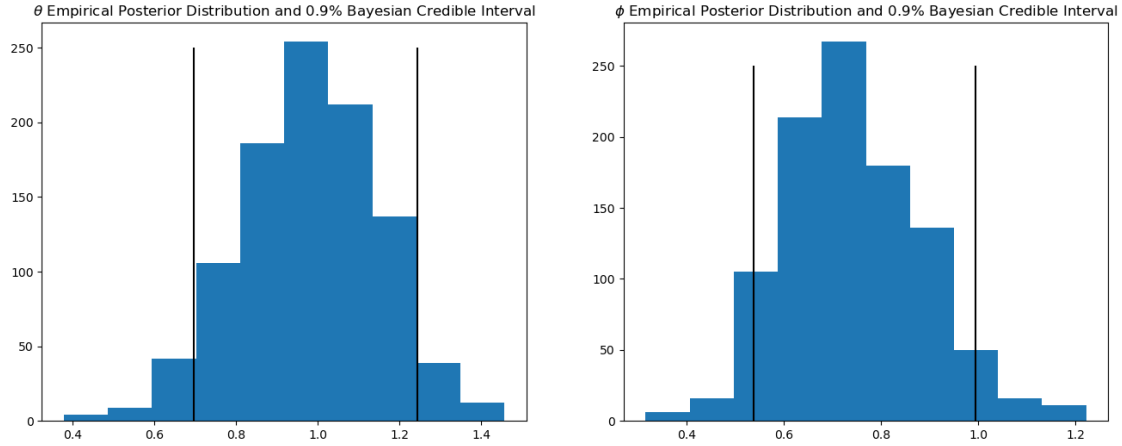
This is because Bayesian posterior inference takes two forms:

1. Estimating the *posterior mean* (expected value integration problem) via *Markov Chain Monte Carlo Integration*
2. Solving the interval bounds of a (probability calculation) integration problem giving a “plausible parameter range” probability statements via *Markov Chain Monte Carlo Integration*

Gibbs sampling is not a *chain rule* decomposition of *full (joint) posterior distributions* $p(\theta|x)$ but nonetheless creates (sequentially [*Markov*] dependent) *Markov chain* samples from *full (joint) posterior distribution* $p(\theta|x)$. While sequentially dependent, they are still samples from $p(\theta|x)$ providing posterior MCMC integration inference

MCMC Bayesian Inference: Credible Intervals [10 minutes]

```
[39]: # MCMC estimate posterior means / re-emphasize credible intervals from initial
      ↪problem
burn_in = 20; BayesianCredibleInterval_probability = 0.90
left_probability = (1-BayesianCredibleInterval_probability)/2
right_probability = left_probability + BayesianCredibleInterval_probability
fig,ax = plt.subplots(1,2,figsize=(16,6))
param=['$\theta$', '$\phi$']
for i,posterior_samples in enumerate([theta_gibbs_posterior_mcmc_draws,
      ↪phi_gibbs_posterior_mcmc_draws]):
    BayesianCredibleInterval = np.quantile(posterior_samples[burn_in:
      ↪], [left_probability, right_probability])
    ax[i].hist(posterior_samples[20:]); ax[i].set_title(param[i]+" Empirical
      ↪Posterior Distribution and "+str(BayesianCredibleInterval_probability)+"%
      ↪Bayesian Credible Interval")
    ax[i].vlines(BayesianCredibleInterval,ymin=0,ymax=250,color = 'k')
```

PyMC: Probabilistic Programming [10 minutes]

- *Hamiltonian Monte Carlo (HMC)* and *Metropolis-Hastings (M-H)* were mentioned when introducing *Gibbs sampling*...
- Without yet knowing what these are, we can use them with PyMC

```
[38]: theta_0,tau = 0,1; alpha,beta = 2,1/2 # Initialize Prior Hyperparameters:
      ↪ specify prior distributions

import pymc as pm; from scipy import stats; x = stats.norm().rvs(10)
normal_gamma_toy_model = pm.Model()
with normal_gamma_toy_model:
    theta = pm.Normal("theta", mu=0, sigma=1)
    phi = pm.Gamma("phi", alpha=1, beta=1)
    x_obs = pm.Normal("likelihood", mu=theta, sigma=1/phi**0.5, observed=x)
```

which specifies...

$$\begin{aligned}
 p(\theta, \phi | x) &\propto p(\theta, \phi, x) = \frac{N(\theta, \phi) N(\theta_0, \tau) \text{Gamma}(\alpha, \beta)}{p(x|\theta) p(\theta) p(\phi)} \quad (\theta \perp\!\!\!\perp \phi) \leftarrow \text{priors} \\
 &= \left[\prod_{i=1}^n \sqrt{\frac{\phi}{2\pi}} e^{-\frac{\phi(x_i - \theta)^2}{2}} \right] \sqrt{\frac{\tau}{2\pi}} e^{-\frac{\tau(\theta - \theta_0)^2}{2}} \frac{\beta^\alpha}{\Gamma(\alpha)} \phi^{\alpha-1} e^{-\beta\phi}
 \end{aligned}$$

PyMC: Posterior Sampling [5 minutes]

- *Hamiltonian Monte Carlo (HMC)* and *Metropolis-Hastings (M-H)* with PyMC
 - *Gibbs sampling* is not even an option for *continuous variables* in PyMC...

```
[61]: with normal_gamma_toy_model:
      idata_HMC = pm.sample() # default is the NUTS implementation of HMC
```

Auto-assigning NUTS sampler...
 Initializing NUTS using jitter+adapt_diag...
 Multiprocess sampling (4 chains in 4 jobs)
 NUTS: [theta, phi]

100.00% [8000/8000 00:00<00:00 Sampling 4 chains, 0 divergences]

Sampling 4 chains for 1_000 tune and 1_000 draw iterations (4_000 + 4_000 draws total) took 0 seconds.

```
[42]: with normal_gamma_toy_model:
      MH = pm.Metropolis([theta, phi], S=np.array([0.1]), tune=False,
      ↪tune_interval=0)
      idata_MH = pm.sample(step=MH) # force MH
```

Multiprocess sampling (4 chains in 4 jobs)
 CompoundStep
 >Metropolis: [theta]
 >Metropolis: [phi]

100.00% [8000/8000 00:00<00:00 Sampling 4 chains, 0 divergences]

Sampling 4 chains for 0 tune and 2_000 draw iterations (0 + 8_000 draws total) took 0 seconds.

The rhat statistic is larger than 1.01 for some parameters. This indicates problems during sampling. See <https://arxiv.org/abs/1903.08008> for details

The effective sample size per chain is smaller than 100 for some parameters. A higher number is needed for reliable rhat and ess computation. See <https://arxiv.org/abs/1903.08008> for details

PyMC: Traceplot “Convergence” Checks [10 minutes]

Do MCMC chains with different initial values seem to agree?

[Warm-up periods are discarded]

M split chains of N samples created from $M/2$ chains halved in the middle with and Split- \hat{R} statistics are [calculated](#): values of 1 mean indistinguishable chains while values of > 1.05 fail the similarity test

Split- \hat{R} is just an ANOVA style comparison of “within chains” variability versus “between chains” variability


Split- $\hat{R} =$

$$\sqrt{\frac{\frac{N-1}{N}W + \frac{1}{M-1} \sum_{m=1}^M (\bar{\theta}^{(m,\cdot)} - \bar{\theta}^{(\cdot,\cdot)})^2}{\frac{1}{M} \sum_{m=1}^M \frac{1}{N-1} \sum_{n=1}^N (\theta^{(m,n)} - \bar{\theta}^{(m,\cdot)})^2}}$$

W : within chain variance

```
[43]: import arviz as az; import matplotlib.pyplot as plt; fig,ax = plt.subplots(2,4,
      ↪figsize=(14,3))
      az.plot_trace(idata_HMC, axes=ax[:2,:2]); az.plot_trace(idata_MH, axes=ax[:2,2:])
      print("HMC Split-Rhats", {k: np.round(v.values,4) for k,v in az.rhat(idata_HMC).
      ↪items()}), end=' | '); print('MH Split-Rhats',{k: np.round(v.values,4) for k,v
      ↪in az.rhat(idata_MH).items()})
```

HMC Split-Rhats {'theta': 1.0026, 'phi': 1.0007} | MH Split-Rhats {'theta': 1.0129, 'phi': 1.0277}




MCMC: Autocorrelation (10 minutes)

```
[62]: import numpy as np; import matplotlib.pyplot as plt; fig,ax = plt.subplots(1,3,
      ↳figsize=(14,4))
max_mcmc_samples = 1000
mcmc_chain = idata_HMC.posterior["phi"].sel(chain=0, draw=slice(0,
      ↳max_mcmc_samples-2)).values; mcmc_chain_one_step_ahead = idata_HMC.
      ↳posterior["phi"].sel(chain=0, draw=slice(1, max_mcmc_samples-1)).values; ax[0].
      ↳plot(mcmc_chain, mcmc_chain_one_step_ahead, '.', alpha=0.5)
print("Autocorrelation for HMC (MCMC) chain", np.
      ↳corrcoef(mcmc_chain,mcmc_chain_one_step_ahead)[0,1], "          ...is
      ↳negative autocorrelation good?")
mcmc_chain = idata_MH.posterior["phi"].sel(chain=0, draw=slice(0,
      ↳max_mcmc_samples-2)).values; mcmc_chain_one_step_ahead = idata_MH.
      ↳posterior["phi"].sel(chain=0, draw=slice(1, max_mcmc_samples-1)).values; ax[1].
      ↳plot(mcmc_chain, mcmc_chain_one_step_ahead, '.', alpha=0.5);
print("Autocorrelation for MH (MCMC) chain", np.
      ↳corrcoef(mcmc_chain,mcmc_chain_one_step_ahead)[0,1], "          ...is
      ↳large positive autocorrelation good?")
ax[2].plot(phi_gibbs_posterior_mcmc_draws[burn:
      ↳-1],phi_gibbs_posterior_mcmc_draws[(burn+1):], '.', alpha=0.5);
print("Autocorrelation for Gibbs Sampling (MCMC) chain", np.
      ↳corrcoef(phi_gibbs_posterior_mcmc_draws[burn:
      ↳-1],phi_gibbs_posterior_mcmc_draws[(burn+1):])[0,1], "    ...is no
      ↳autocorrelation the best?")
```

```
Autocorrelation for HMC (MCMC) chain -0.0140628109479451          ...is
negative autocorrelation good?
Autocorrelation for MH (MCMC) chain 0.9084368880554985          ...is
```

large positive autocorrelation good?

Autocorrelation for Gibbs Sampling (MCMC) chain 0.021150850107231244 ...is no autocorrelation the best?



MCMC: Effective Sample Size [10 minutes]

PyMC labels $n_{\text{eff}} = \frac{n}{\sum_{t=-\infty}^{\infty} \rho_t} = \frac{n}{1+2\sum_{t=1}^{\infty} \rho_t}$ for order t autocorrelations ρ_t in the MCMC chains

`ess_bulk` while `ess_tail` is the same but after converting to binary sequences estimating 5% tail quantiles.

```
[69]: import arviz as az; print("number of MCMC samples:", idata_HMC.posterior['phi'].
      ↪ values.flatten().shape[-1])
      az.summary(idata_HMC, round_to=2) # https://python.arviz.org/en/stable/api/
      ↪ generated/arviz.ess.html
```

number of MCMC samples: 4000

```
[69]:
```

	mean	sd	hdi_3%	hdi_97%	mcse_mean	mcse_sd	ess_bulk	ess_tail	\
theta	0.98	0.17	0.66	1.31	0.0	0.0	3880.51	2568.12	
phi	0.70	0.14	0.45	0.96	0.0	0.0	3959.53	2736.66	


```

      r_hat
theta    1.0
phi      1.0
```

```
[68]: print("number of MCMC samples:", idata_MH.posterior['phi'].values.flatten().
      ↪ shape[-1])
      az.summary(idata_MH, round_to=2)
```

number of MCMC samples: 8000

```
[68]:
```

	mean	sd	hdi_3%	hdi_97%	mcse_mean	mcse_sd	ess_bulk	ess_tail	\
theta	0.97	0.18	0.65	1.30	0.01	0.01	390.50	561.12	
phi	0.69	0.14	0.42	0.97	0.01	0.01	188.56	317.92	

```

r_hat
theta  1.01
phi    1.03

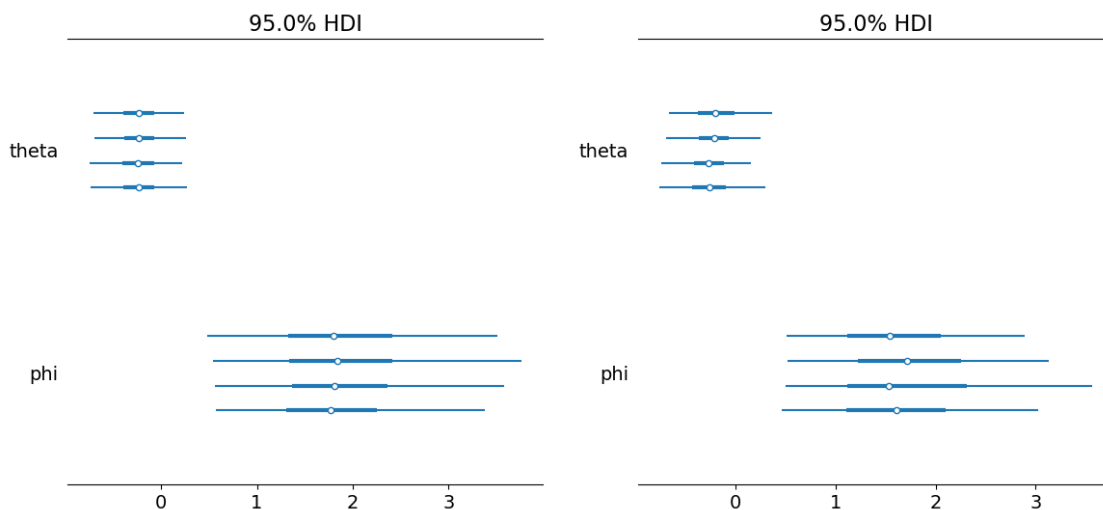
```

MCMC Bayesian Inference: Credible Intervals [5 minutes]

```

[258]: import arviz as az; import matplotlib.pyplot as plt; fig,ax = plt.subplots(1,2,
↳figsize=(14,6))
az.plot_forest(idata_HMC, var_names=["theta","phi"], combined=False, hdi_prob=0.
↳95, ax=ax[0]);
az.plot_forest(idata_MH, var_names=["theta","phi"], combined=False, hdi_prob=0.
↳95, ax=ax[1]);

```



Homework #3 – *you'll need to do this on google colab*

- Find data set on [kaggle](#).
- Provide posterior inference for a column of data with a PyMC *normal-gamma specification*.
- Find a different column of data for which a *normal-gamma specification* isn't justified.
- Create a different PyMC specification that IS appropriate for this new data column:
 - here are your PyMC choices for [continuous distributions](#)
 - here are your PyMC choices for [discrete distributions](#)

...just make sure you've chosen a reasonable likelihood for your non-normal data and that the support of your priors makes sense for the values the parameters of your likelihood can take on...

- Provide posterior inference for the parameters of your new specification given the column of non-normal data using PyMC and provide an analysis of the MCMC chains.