

Class Activity 6: Quiz [10 minutes]

1. What's the difference between these three data models. Can they be used for Bayesian inference? How?

$$\begin{aligned}
 \mathbf{y}_{n \times 1} &\sim \mathcal{MVN}(\mathbf{X}_{n \times p} \boldsymbol{\beta}_{p \times 1}, \boldsymbol{\Sigma}_{n \times n} = \sigma^2 \mathbf{I}_{n \times n}) \\
 p(\mathbf{y} | \mathbf{X}, \boldsymbol{\beta}, \Sigma) &= (2\pi\sigma^2)^{-n/2} \exp\left(-\frac{1}{2\sigma^2} (\mathbf{y} - \mathbf{X}\boldsymbol{\beta})^\top (\mathbf{y} - \mathbf{X}\boldsymbol{\beta})\right) \\
 \hline
 \mathbf{y}_i &\sim \mathcal{MVN}(\boldsymbol{\mu}_{p \times 1}, \boldsymbol{\Sigma}_{p \times p} = \mathbf{D}\mathbf{R}\mathbf{D}) \\
 p(\mathbf{y} | \boldsymbol{\mu}, \Sigma) &= (2\pi)^{-n/2} \det(\boldsymbol{\Sigma})^{-1/2} \exp\left(-\frac{1}{2} (\mathbf{y} - \boldsymbol{\mu})^\top \boldsymbol{\Sigma}^{-1} (\mathbf{y} - \boldsymbol{\mu})\right) \\
 \hline
 y_i &\sim \text{Bernoulli}(\Pr(y_i = 1 | \mathbf{x}_i) = \Phi(\mathbf{x}_i^\top \boldsymbol{\beta}_{p \times 1})) \quad \text{Hint: or use } \Pr(y_i = 1 | \mathbf{x}_i) = \frac{1}{1 + e^{-\mathbf{x}_i^\top \boldsymbol{\beta}}} \\
 &\quad \text{CDF of } \mathcal{N}(0,1) \\
 p(y | \mathbf{x}, \boldsymbol{\beta}) &= \Phi(\mathbf{x}^\top \boldsymbol{\beta}_{p \times 1})^y (1 - \Phi(\mathbf{x}^\top \boldsymbol{\beta}_{p \times 1}))^{1-y}
 \end{aligned}$$

2. How could the middle model be used to create **multivariate linear regression** predicting a **vector** outcome \mathbf{y} (as opposed to just **linear regression** predicting a **univariate** y_i , which is actually what the first model does...)?
3. Show that $\Phi(\mathbf{x}^\top \boldsymbol{\beta}_{p \times 1}) = \Pr(z \geq 0)$ for $z \sim \mathcal{N}(\mathbf{x}^\top \boldsymbol{\beta}_{p \times 1}, 1)$
 Hint: show that $\int_{-\infty}^{\mathbf{x}^\top \boldsymbol{\beta}_{p \times 1}} e^{-\frac{1}{2}z^2} dz = \int_0^{\infty} e^{-\frac{1}{2}(z - \mathbf{x}^\top \boldsymbol{\beta}_{p \times 1})^2} dz$

Class Activity 6: Solutions + Another Question [10 minutes]

1. The first is **linear regression model** where the features \mathbf{X} are used to predict outcomes \mathbf{y} ; whereas, the second is simply a **multivariate normal distribution**; whereas, the third is a **generalized linear model** predicting a binary outcome (often called a **classification model**) parameterizing the chance of "success" using either a **probit** $\Phi(\mathbf{x}^\top \boldsymbol{\beta}_{p \times 1})$ or **logit** $\frac{1}{1 + e^{-\mathbf{x}^\top \boldsymbol{\beta}}}$ ***link function**.

For Bayesian inference on β , σ^2 or $\boldsymbol{\mu}$ or $\boldsymbol{\Sigma}$ or β , just put priors on these parameters...

2. Each element $\mu_{ik} = \mathbf{x}_i^T \boldsymbol{\beta}_{(k)}$ of $\boldsymbol{\mu}_i$ becomes a **linear model** and then the **covariance matrix** $\boldsymbol{\Sigma}$ captures the **residual covariance dependence structure** observed over the **residual multivariate outcomes** ($\mathbf{y}_i - \boldsymbol{\mu}_i$)

$$\mu_{ik} = \mathbf{x}_i^T \boldsymbol{\beta}_{(k)} \quad \text{and} \quad p(\mathbf{y}|\boldsymbol{\mu}, \boldsymbol{\Sigma}) = (2\pi)^{-n/2} \det(\boldsymbol{\Sigma})^{-1/2} \exp\left(-\frac{1}{2}(\mathbf{y}_i - \boldsymbol{\mu}_i)^\top \boldsymbol{\Sigma}^{-1}(\mathbf{y}_i - \boldsymbol{\mu}_i)\right)$$

3. $\Phi(\mathbf{x}_i^\top \boldsymbol{\beta}_{p \times 1}) = \Pr(z_i \geq 0)$ for $z_i \sim \mathcal{N}(\mathbf{x}_i^\top \boldsymbol{\beta}_{p \times 1}, 1)$ since

$$\int_{-\infty}^{\mathbf{x}_i^\top \boldsymbol{\beta}_{p \times 1}} e^{-\frac{1}{2}z_i^2} dz = \int_{-\infty}^0 e^{-\frac{1}{2}(z_i + \mathbf{x}_i^\top \boldsymbol{\beta}_{p \times 1})^2} dz = \int_{-\infty}^0 e^{-\frac{1}{2}(-z_i - \mathbf{x}_i^\top \boldsymbol{\beta}_{p \times 1})^2} dz = \int_0^\infty e^{-\frac{1}{2}(z_i - \mathbf{x}_i^\top \boldsymbol{\beta}_{p \times 1})^2} dz$$

4. Specify a **multivariate generalized linear regression** model whose **marginal distributions** are each univariate **bernoulli distributions** but which also **parameterizes a joint dependency** structure between **marginal outcomes** Hint:
 $\Pr(y_{ij} = 1|\mathbf{x}_i) = \Pr(z_{ij} \geq 0)$ and specify a joint distribution for unobserved \mathbf{z}_i that depends on \mathbf{x}_i

Class Activity 6: Solutions [5 minutes]

4. Specify a **multivariate generalized linear regression** model whose **marginal distributions** are each univariate **bernoulli distributions** but which also **parameterizes a joint dependency** structure between **marginal outcomes** Hint:
 $\Pr(y_{ij} = 1|\mathbf{x}_i) = \Pr(z_{ij} \geq 0)$ and specify a joint distribution for unobserved \mathbf{z}_i that depends on \mathbf{x}_i

Generalized Multivariate Linear Model (GMLM)

- $\Pr(y_{ij} = 1|\mathbf{x}_i) = \Pr(z_{ij} \geq 0)$
- $\mathbf{z}_i \sim \mathcal{MVN}(\boldsymbol{\mu}_i, \boldsymbol{\Sigma} = \mathbf{D}\mathbf{R}\mathbf{D} = \mathbf{R})$

$$p(\mathbf{z}_i|\boldsymbol{\mu}_i, \boldsymbol{\Sigma} = \mathbf{D}\mathbf{R}\mathbf{D} = \mathbf{R}) = (2\pi)^{-n/2} \det(\mathbf{R})^{-1/2} \exp\left(-\frac{1}{2}(\mathbf{z}_i - \boldsymbol{\mu}_i)^\top \mathbf{R}^{-1}(\mathbf{z}_i - \boldsymbol{\mu}_i)\right)$$

- with $\mu_{ij} = \mathbf{x}_i^T \boldsymbol{\beta}_{(j)}$

and where \mathbf{R} is used in the same spirit as $\Pr(y_{ij} = 1|\mathbf{x}_i) = \Phi(\mathbf{x}_i^\top \boldsymbol{\beta}_{p \times 1}) = \Pr(z_i \geq 0)$ for $z_i \sim \mathcal{N}(\mathbf{x}_i^\top \boldsymbol{\beta}_{p \times 1}, 1)$ but additionally captures **residual correlation dependence structure** observed over the **residual multivariate outcomes**

The transformation on the **marginal distribution** to the **probability parameter** $\Pr(y_{ij} = 1|\mathbf{x}_i)$ of a **Bernoulli distribution** could be carried out with the **logit** $\frac{1}{1+e^{-z_{ij}}}$ **link function** instead of the **probit** $\Phi(\mu_{ij} = \mathbf{x}_i^\top \boldsymbol{\beta}_{(j)})$ **link function** as above

- The transformation of the **marginal distributions** of a **multivariate normal random variable** into other **distributional forms** (such as **Bernoulli distributions** as done here) is the core principle behind **copulas**

Homework 6: Part I

1. Go get data from kaggle.com and do a **(Univariate) Bayesian Logistic Regression** analysis
2. Adjust the code below to specify that the outcomes have a Bernoulli distribution and use a **logit** or **probit link function** (or $\Pr(z \leq 0)$ for latent z) to correctly parameterize the predicted values of the observed outcomes

```
import pymc as pm; import numpy as np
n,p=100,10; X,y=np.zeros((n,p)),np.ones((n,1))
# Replace this made up data with your data set from kaggle...
with pm.Model() as MLR:
    betas = pm.MvNormal('betas', mu=np.zeros((p,1)), cov=np.eye(p), shape=(p,1))
    sigma = pm.TruncatedNormal('sigma', mu=1, sigma=1, lower=0) # half normal
    y = pm.Normal('y', mu=pm.math.dot(X, betas), sigma=sigma, observed=y)

with MLR:
    idata = pm.sample()
```

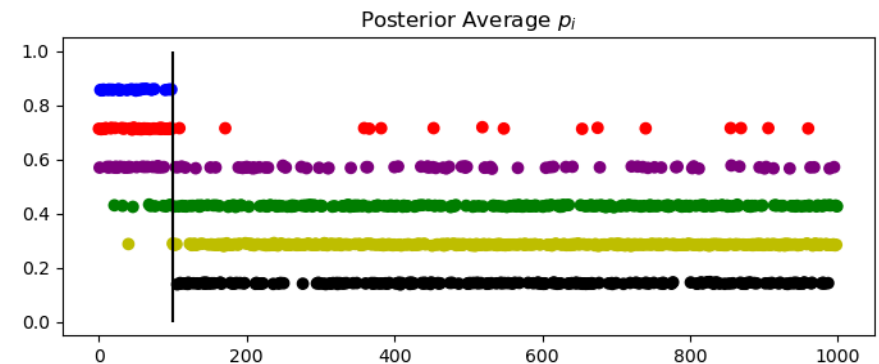
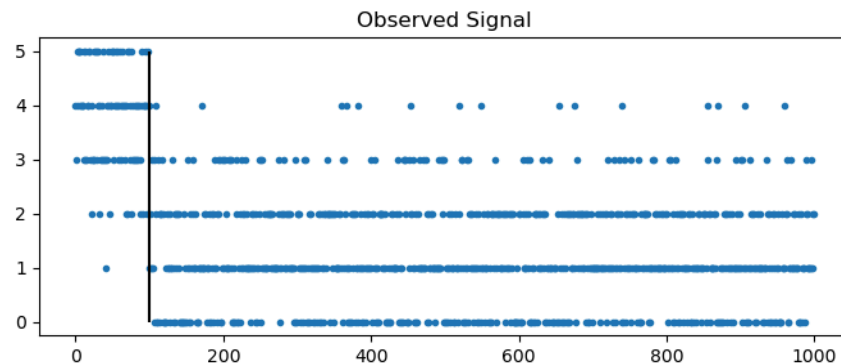
3. Choose **prior** that are sensible for your specification
4. [Optional] Assess the performance of the MCMC and any issues or warnings in the **standard manner**
5. [Optional] Go get data from kaggle.com and do a **Multivariate Bayesian Logistic Regression** analysis

Bayesian Multiplicity Adjustment: Significance Dichotomy [7 minutes]

Hypothesis testing involves a mixture of **effect sizes**, some large and some negligible

This can coarsely be viewed as the dichotomy of **significant** versus **insignificant tests**

```
In [ ]: import numpy as np; from scipy import stats; import matplotlib.pyplot as plt;
fig,ax = plt.subplots(1,2,figsize=(18,3))
# first 100 are "interesting" (high p) distributions;
# but the (posterior p) "evidence" depends only on observed outcome
np.random.seed(1);
y_obs = np.r_[stats.binom(p=0.75,n=5).rvs(100),stats.binom(p=0.25,n=5).rvs(900)];
ax[0].plot(y_obs,'.'); ax[0].set_title('Observed Signal');
ax[0].vlines(100,ymin=0,ymax=5,color='k');
ax[1].scatter(np.arange(1000), idata.posterior['p'].values.reshape((-1,1000)).mean(axis=0),c=[['k','y','g','purple']
ax[1].set_title('Posterior Average $p_i$');
ax[1].vlines(100,ymin=0,ymax=1,color='k');
```



```
In [ ]: import pymc as pm
with pm.Model() as no_multiplicity_correction:
    p = pm.Beta('p', alpha=1, beta=1, shape=1000)
    y = pm.Binomial('y', n=5, p=p, observed=y_obs)
```

```
with no_multiplicity_correction:
    idata = pm.sample()
```

Auto-assigning NUTS sampler...

Initializing NUTS using jitter+adapt_diag...

Multiprocess sampling (4 chains in 4 jobs)

NUTS: [p]

100.00% [8000/8000 00:03<00:00 Sampling 4 chains, 0 divergences]

Sampling 4 chains for 1_000 tune and 1_000 draw iterations (4_000 + 4_000 draws total) took 3 seconds.

Bayesian Multiplicity Adjustment: Hierarchical Modeling [10 minutes]

$$\begin{array}{lll}
 \cancel{p_i} \sim \cancel{\text{beta}(\alpha, \beta)} & p \sim \text{beta}(\alpha, \beta) & x_i \sim \text{Bernoulli}(p) & p_0 \sim \text{beta}(\alpha_0, \beta_0) \\
 \cancel{y_i} \sim \cancel{\text{binomial}(p, n)} & y_i \sim \text{binomial}(\underbrace{p_0 + x_i p_1}_{\text{restrict to } [0,1]}, n) & & p_1 \sim \text{beta}(\alpha_1, \beta_1)
 \end{array}$$

```
In [ ]: with pm.Model() as multiplicity_correction:
    p = pm.Beta('p', alpha=1, beta=1)
    x = pm.Binomial('x', n=1, p=p, shape=1000)
    p0 = pm.Beta('p0', alpha=1, beta=1)
    p1 = pm.Beta('p1', alpha=1, beta=1)
    no_negatives = pm.math.switch(pm.math.lt(p0+x*p1,0), 0, p0+x*p1)
    in_unit_interval = pm.math.switch(pm.math.gt(no_negatives,1), 1, p0+x*p1)
    p01 = pm.Deterministic('p01', in_unit_interval)
    y = pm.Binomial('y', n=5, p=p01, observed=y_obs) # p=0.25+x*0.5

    with multiplicity_correction:
        idata2 = pm.sample()
```

Multiprocess sampling (4 chains in 4 jobs)

CompoundStep

>NUTS: [p, p0, p1]

>Metropolis: [x]

100.00% [8000/8000 02:34<00:00 Sampling 4 chains, 0 divergences]

Sampling 4 chains for 1_000 tune and 1_000 draw iterations (4_000 + 4_000 draws total) took 155 seconds.

/Users/scottschwartz/miniconda3/envs/PyMC/lib/python3.11/site-packages/arviz/stats/diagnostics.py:592: RuntimeWarning: invalid value encountered in scalar divide

(between_chain_variance / within_chain_variance + num_samples - 1) / (num_samples)

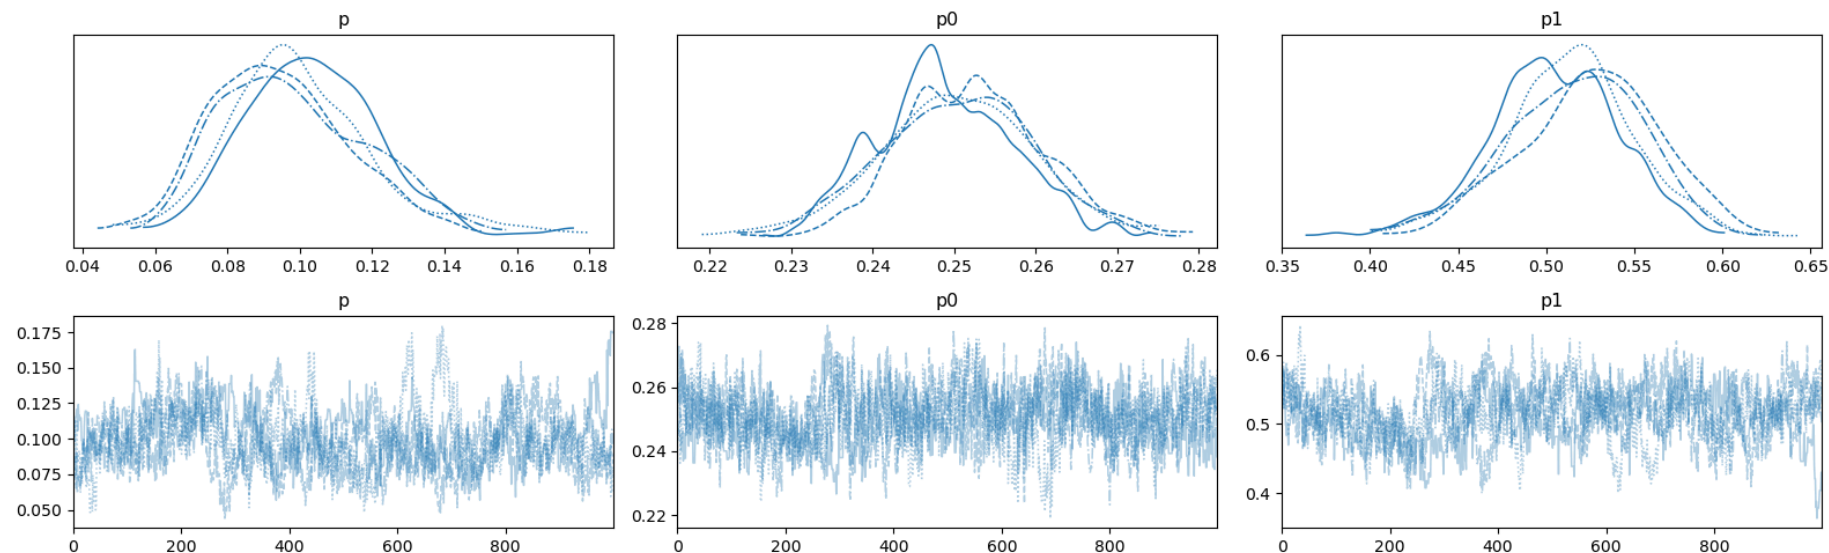
The rhat statistic is larger than 1.01 for some parameters. This indicates problems during sampling. See <https://arxiv.org/abs/1903.08008> for details

The effective sample size per chain is smaller than 100 for some parameters. A higher number is needed for reliable rhat and ess computation. See <https://arxiv.org/abs/1903.08008> for details

Bayesian Multiplicity Adjustment: Hierarchical Modeling [3 minutes]

$$\begin{array}{lll}
 \cancel{p_i} \sim \cancel{\text{beta}(\alpha, \beta)} & p \sim \text{beta}(\alpha, \beta) & x_i \sim \text{Bernoulli}(p) & p_0 \sim \text{beta}(\alpha_0, \beta_0) \\
 \cancel{y_i} \sim \cancel{\text{binomial}(p, n)} & y_i \sim \text{binomial}(\underbrace{p_0 + x_i p_1}_{\text{restrict to } [0,1]}, n) & & p_1 \sim \text{beta}(\alpha_1, \beta_1)
 \end{array}$$

```
In [ ]: fig, ax = plt.subplots(2, 3, figsize=(16, 5))
import arviz as az; az.plot_trace(idata2, var_names=["p", "p0", "p1"], axes=ax.T);
plt.tight_layout();
```

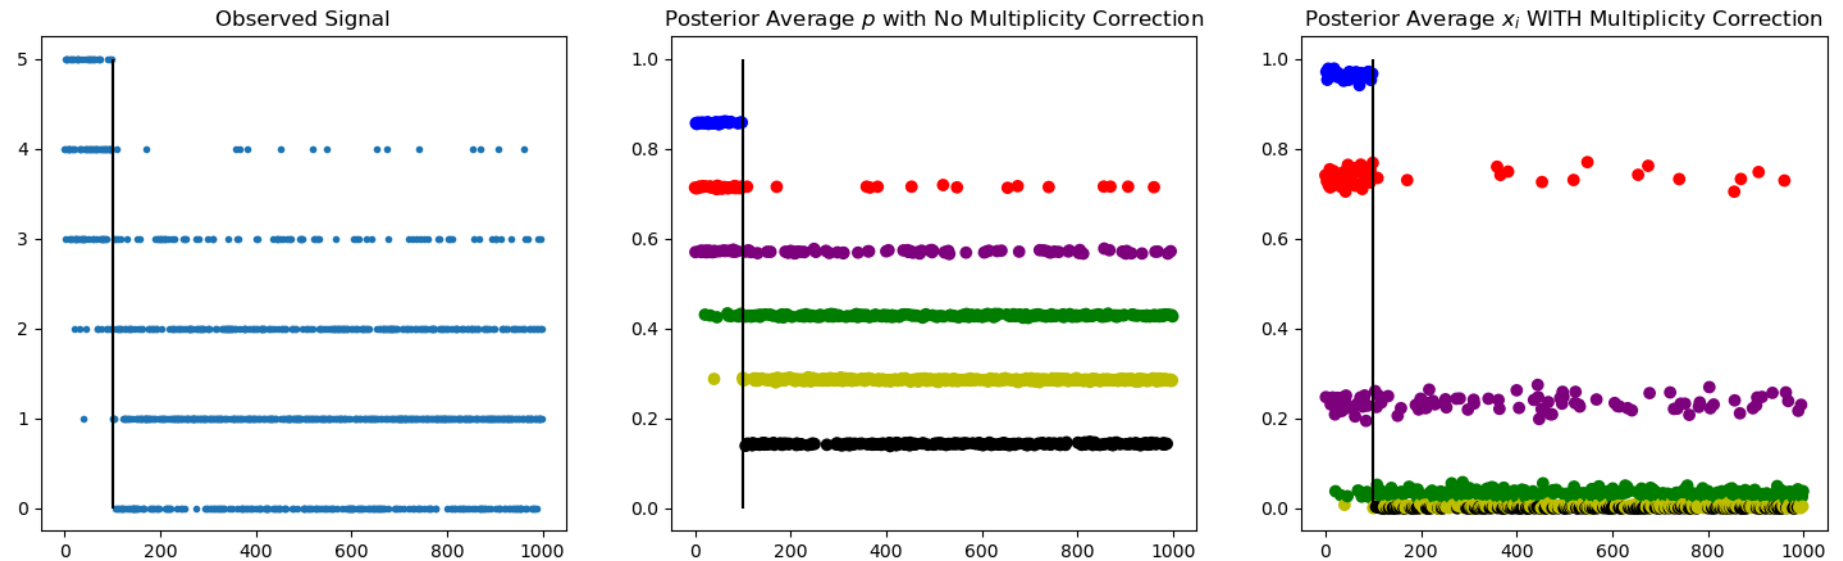


Bayesian Multiplicity Adjustment: Hierarchical Modeling and the Power of Information Sharing [5 minutes]

Learning the proportion of examples that are "interesting" actually means learning

WHICH examples are interesting by **imputing** weather or not each of the examples is interesting one or not

```
In [ ]: # first 100 are "interesting" (high p) distributions; but the (posterior p) "evidence" depends only on observed out
fig, ax = plt.subplots(1, 3, figsize=(18, 5)); np.random.seed(1);
y_obs = np.r_[stats.binom(p=0.75, n=5).rvs(100), stats.binom(p=0.25, n=5).rvs(900)];
ax[0].plot(y_obs, '.'); ax[0].set_title('Observed Signal');
ax[0].vlines(100, ymin=0, ymax=5, color='k');
ax[1].scatter(np.arange(1000), idata.posterior['p'].values.reshape((-1, 1000)).mean(axis=0),
              c=[['k', 'y', 'g', 'purple', 'r', 'b'][c] for c in y_obs]);
ax[1].set_title('Posterior Average $p$ with No Multiplicity Correction');
ax[1].vlines(100, ymin=0, ymax=1, color='k');
ax[2].scatter(np.arange(1000), idata2.posterior['x'].values.reshape((-1, 1000)).mean(axis=0),
              c=[['k', 'y', 'g', 'purple', 'r', 'b'][c] for c in y_obs]);
ax[2].set_title('Posterior Average $x_{i|}$ WITH Multiplicity Correction');
ax[2].vlines(100, ymin=0, ymax=1, color='k');
```



Bayesian Variable Selection: "Spike and Slab" [12 minutes]

$$\begin{aligned}
 p &\sim \text{beta}(\alpha, \beta) \\
 s_i &\sim \text{Bernoulli}(p) & b_i &\sim \text{Normal}(\mu_0, \sigma_0) & \beta_i &= b_i \times s_i \\
 y_i &\sim \text{Normal}(x_i^T \beta, \sigma) & \sigma &\sim \text{HalfNormal}(\sigma_0)
 \end{aligned}$$

```

In [ ]: m,q = 20,10; betas = np.zeros((m,1)); betas[0:q,0] = np.linspace(0,q-1,q); np.random.seed(2)
n = 100; X = stats.binom(n=1,p=0.5).rvs(size=(n,m)); y_obs=X.dot(betas).flatten() + stats.norm().rvs(size=n)

with pm.Model() as spikeNslab:
    p = pm.Beta('p', alpha=1, beta=1)
    spike = pm.Binomial('spike', n=1, p=p, shape=m)
    slab = pm.Normal('slab', mu=0, sigma=10, shape=m)
    beta = pm.Deterministic('beta', spike*slab) # elementwise multiplication
    sigma = pm.HalfNormal('sigma', sigma=2)
    y = pm.Normal('y', mu=pm.math.dot(X, beta), sigma=sigma, observed=y_obs)

```



```
with spikeNslab:
    idata3 = pm.sample()
```

Multiprocess sampling (4 chains in 4 jobs)

CompoundStep

>NUTS: [p, slab, sigma]

>Metropolis: [spike]

100.00% [8000/8000 00:07<00:00 Sampling 4 chains, 1,033 divergences]

Sampling 4 chains for 1_000 tune and 1_000 draw iterations (4_000 + 4_000 draws total) took 8 seconds.

/Users/scottschwartz/miniconda3/envs/PyMC/lib/python3.11/site-packages/arviz/stats/diagnostics.py:592: RuntimeWarning: invalid value encountered in scalar divide

(between_chain_variance / within_chain_variance + num_samples - 1) / (num_samples)

The rhat statistic is larger than 1.01 for some parameters. This indicates problems during sampling. See <https://arxiv.org/abs/1903.08008> for details

The effective sample size per chain is smaller than 100 for some parameters. A higher number is needed for reliable rhat and ess computation. See <https://arxiv.org/abs/1903.08008> for details

There were 1033 divergences after tuning. Increase `target_accept` or reparameterize.

Bayesian Variable Selection: "Spike and Slab" [4 minutes]

$$p \sim \text{beta}(\alpha, \beta)$$

$$s_i \sim \text{Bernoulli}(p)$$

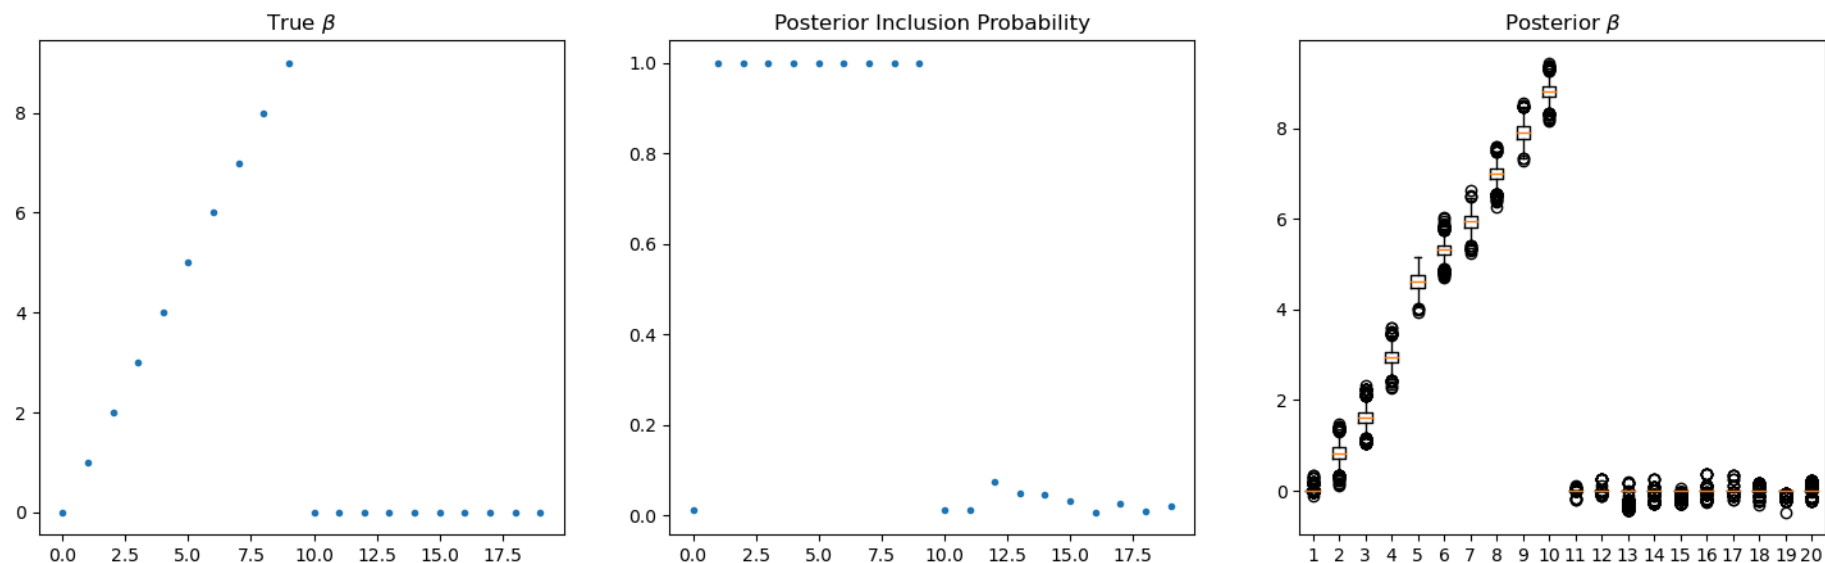
$$y_i \sim \text{Normal}(x_i^T \beta, \sigma)$$

$$b_i \sim \text{Normal}(\mu_0, \sigma_0)$$

$$\beta_i = b_i \times s_i$$

$$\sigma \sim \text{HalfNormal}(\sigma_0)$$

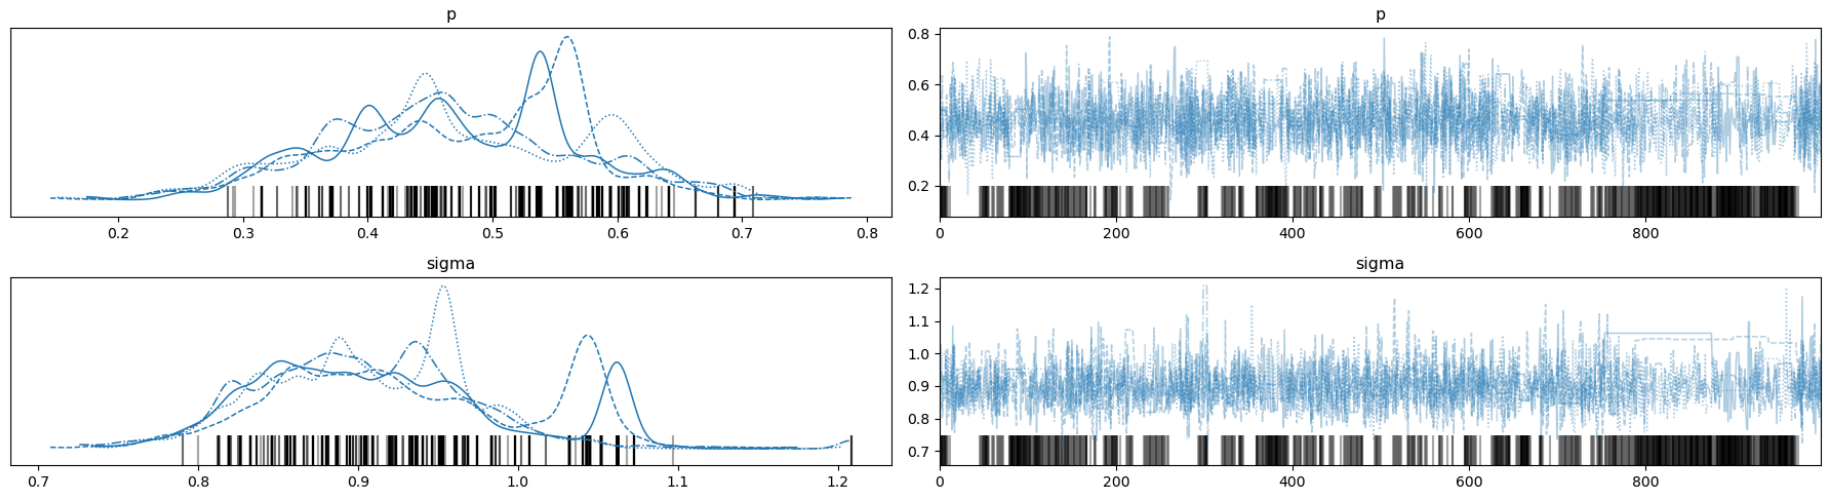
```
In [ ]: fig, ax = plt.subplots(1, 3, figsize=(18, 5)); ax[0].plot(betas, '.'); ax[0].set_title('True $\beta$'); ax[1].plot(idat
```



Bayesian Variable Selection: "Spike and Slab" [4 minutes]

$$\begin{aligned}
 p &\sim \text{beta}(\alpha, \beta) \\
 s_i &\sim \text{Bernoulli}(p) & b_i &\sim \text{Normal}(\mu_0, \sigma_0) & \beta_i &= b_i \times s_i \\
 y_i &\sim \text{Normal}(x_i^T \beta, \sigma) & \sigma &\sim \text{HalfNormal}(\sigma_0)
 \end{aligned}$$

```
In [ ]: import arviz as az; az.plot_trace(idata3, var_names=["p", "sigma"], figsize=(18,5)); plt.tight_layout();
```



Bayesian Shrinkage Estimation: "Lasso Regression" [10 minutes]

Ridge Regression L_2 -penalty is just using *norml priors* for the *regression coefficients*

$$\beta_i \sim \text{Normal}(b_i, s_i) \quad f(\beta_i | b_i, s_i) = \frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{1}{2} \left(\frac{\beta_i - b_i}{s_i} \right)^2}$$

$$y_i \sim \text{Normal}(x_i^T \beta, \sigma) \quad \sigma \sim \text{HalfNormal}(\sigma_0)$$

$$\overbrace{|\beta_i - b_i|^2}^{L_2}$$

Squared Penalization

Lasso Regression L_1 -penalty is just replacing the *norml priors* with **Laplace ("Double Exponential") prior distributions**

$$\beta_i \sim \text{Laplace}(b_i, s_i) \quad f(\beta_i | b_i, s_i) = \frac{1}{2b} \exp\left(-\frac{|\beta_i - b_i|}{s_i}\right)$$

$$y_i \sim \text{Normal}(x_i^T \beta, \sigma) \quad \sigma \sim \text{HalfNormal}(\sigma_0)$$

$$\overbrace{|\beta_i - b_i|}^{L_1}$$

Absolute Penalization

Homework 6: Part II

Regularized Loss Functions

Machine Learning fits models by optimizing penalized **loss functions**

Two classic regularizations are "ridge" and "lasso" regression, which respectively use L_2 and L_1 penalty functions

- Lasso:

$$\sum_{i=1}^n \frac{1}{2} (y_i - x_i^T \beta_{p \times 1})^2 + \lambda \sum_{j=1}^n \beta_j^2 = \frac{1}{2} (y - X\beta)^T (y - X\beta) + \lambda \sum_{j=1}^n \beta_j^2 = \frac{1}{2} \|y - X\beta\|_2^2 + \lambda \|\beta\|_2^2$$

- Ridge:

$$\sum_{i=1}^n \frac{1}{2} (y_i - x_i^T \beta_{p \times 1})^2 + \lambda \sum_{j=1}^n |\beta_j| = \frac{1}{2} (y - X\beta)^T (y - X\beta) + \lambda \sum_{j=1}^n |\beta_j| = \frac{1}{2} \|y - X\beta\|_2^2 + \lambda \|\beta\|_1$$

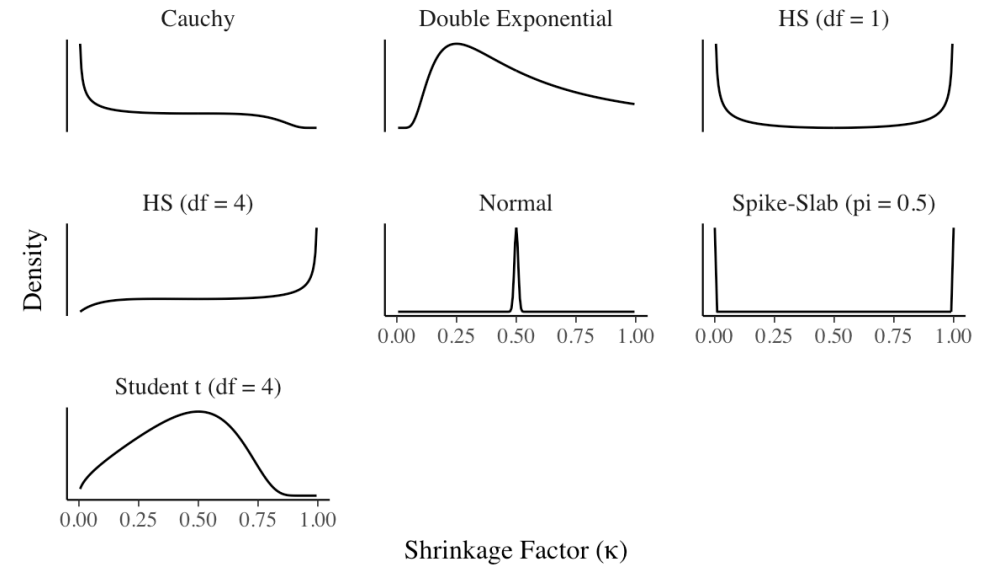
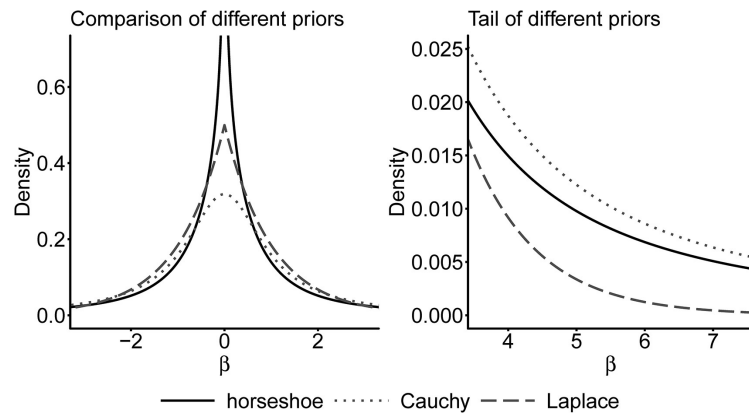
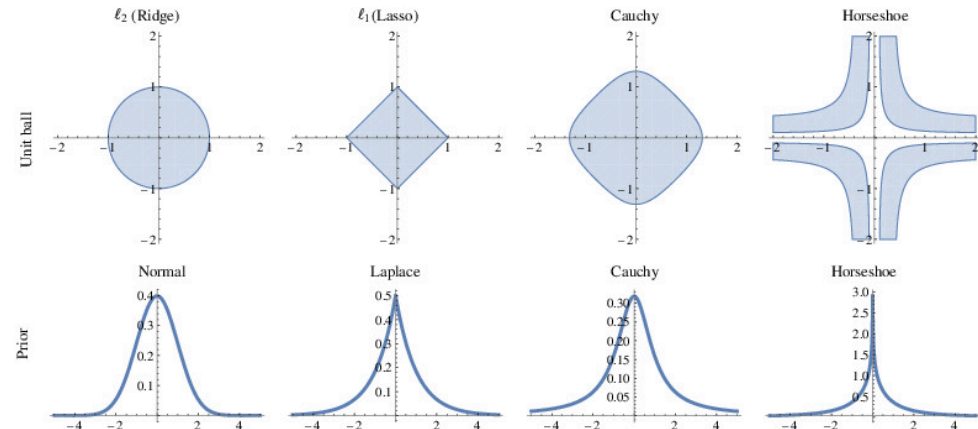
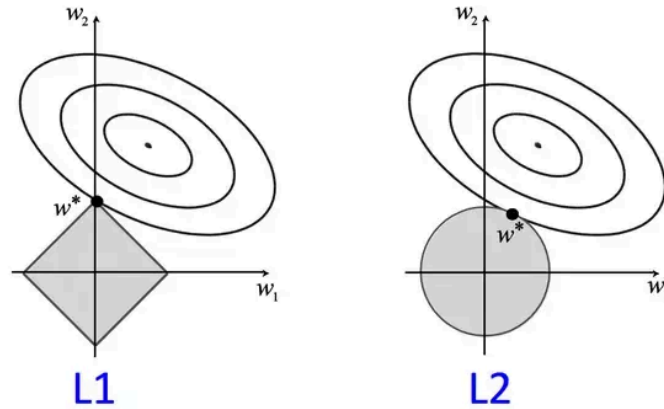
Show that for $\sigma = 1$ and **hyperparameters** $b_i = 0$ (ignoring normalizing proportionality constants) the log posterior distributions for β using either **normal** or **Laplace** prior distributions have analogous forms to the above expressions

Now write down and understand the following: "Bayesians do not optimize posterior distributions, they sample from them; but, the posterior distributions are nonetheless 'regularizations' of the likelihood through the prior."

Bayesian Shrinkage Estimation: Geometrically [15 minutes]

Machine Learning thought they invented "regularization"... *Priors are Regularization!*

But Bayes been regularizing since 1763: Bayesian analysis is the original regularization methodology



Bayesian Shrinkage Estimation: "Lasso Regression" [2 minutes]

```
In [ ]: m,q = 20,10; betas = np.zeros((m,1)); betas[0:q,0] = np.linspace(2,q+1,q); np.random.seed(2)
n = 1000; X = stats.binom(n=1,p=0.5).rvs(size=(n,m)); y_obs=X.dot(betas).flatten() + stats.norm().rvs(size=n)
with pm.Model() as lasso:
    beta = pm.Laplace('beta', mu=0, b=1, shape=m); beta0 = pm.Normal('beta0', mu=0, sigma=10); sigma = pm.HalfNormal
    y = pm.Normal('y', mu=beta0+pm.math.dot(X, beta), sigma=sigma, observed=y_obs)
    idata4 = pm.sample()
```

Auto-assigning NUTS sampler...

Initializing NUTS using jitter+adapt_diag...

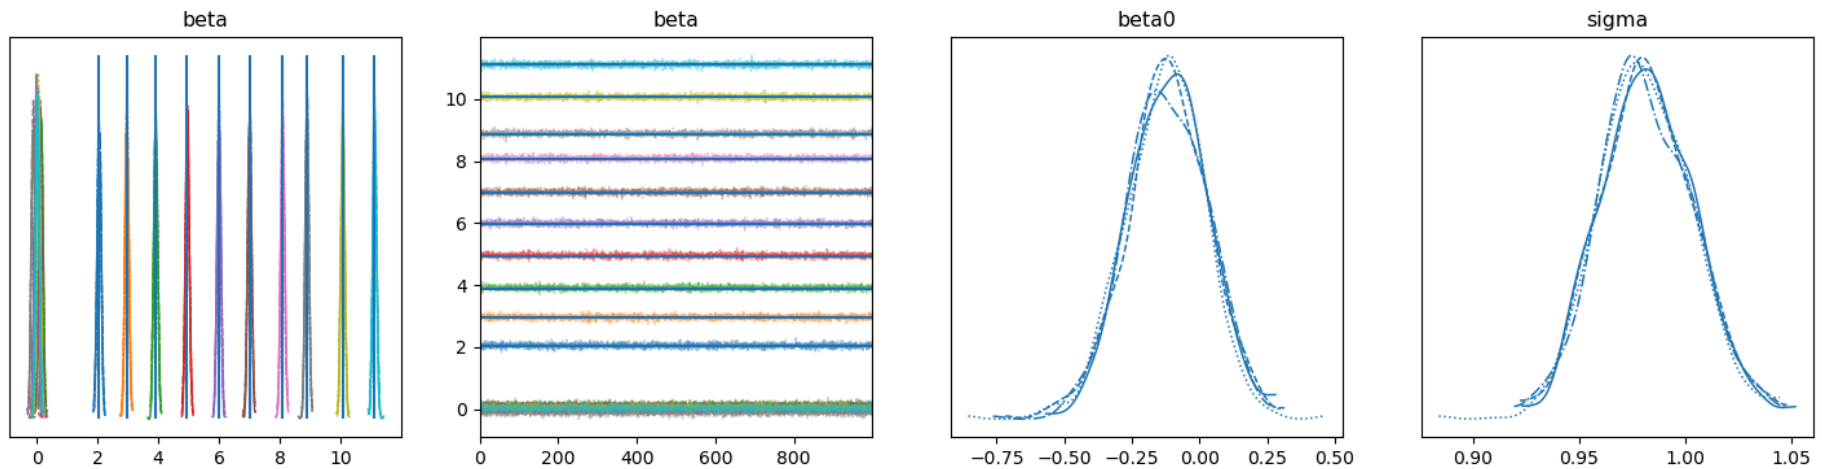
Multiprocess sampling (4 chains in 4 jobs)

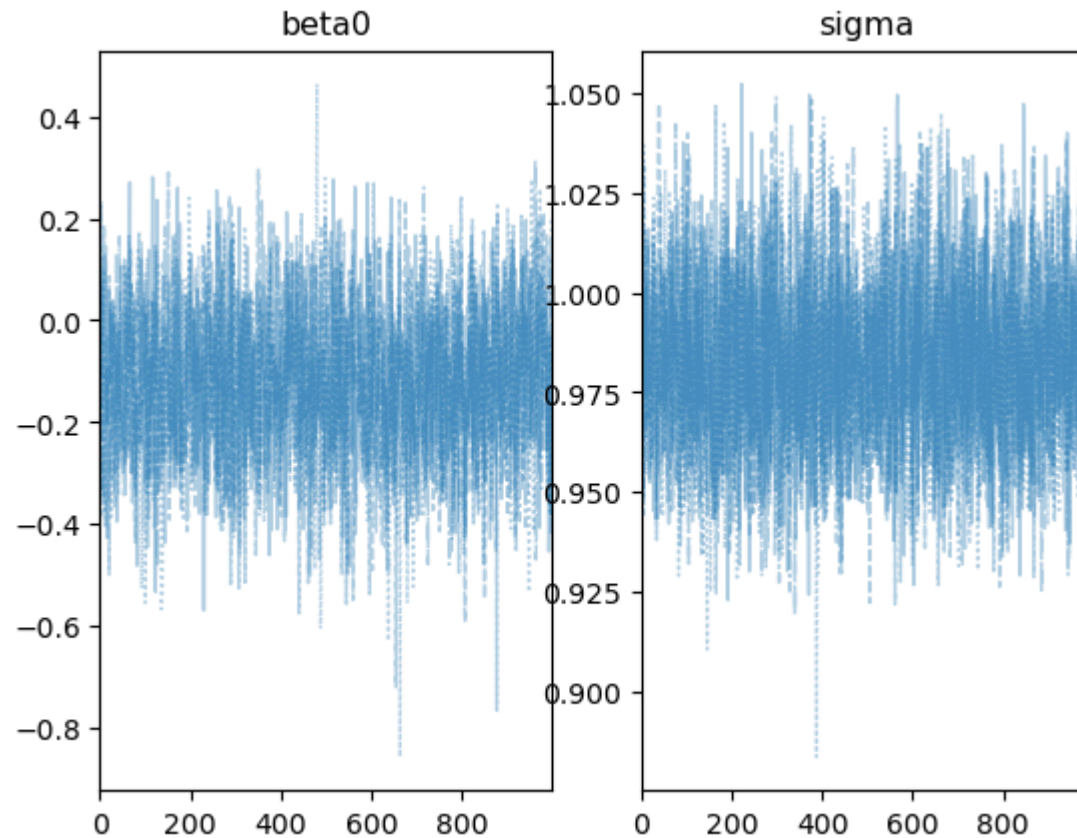
NUTS: [beta, beta0, sigma]

100.00% [8000/8000 00:13<00:00 Sampling 4 chains, 0 divergences]

Sampling 4 chains for 1_000 tune and 1_000 draw iterations (4_000 + 4_000 draws total) took 13 seconds.

```
In [ ]: fig,ax = plt.subplots(1,4,figsize=(18,4)); fig2,ax2 = plt.subplots(1,2); az.plot_trace(idata4, axes=np.r_[ax[:3],ax
```





Bayesian Shrinkage Estimation: "Lasso Regression" [3 minutes]

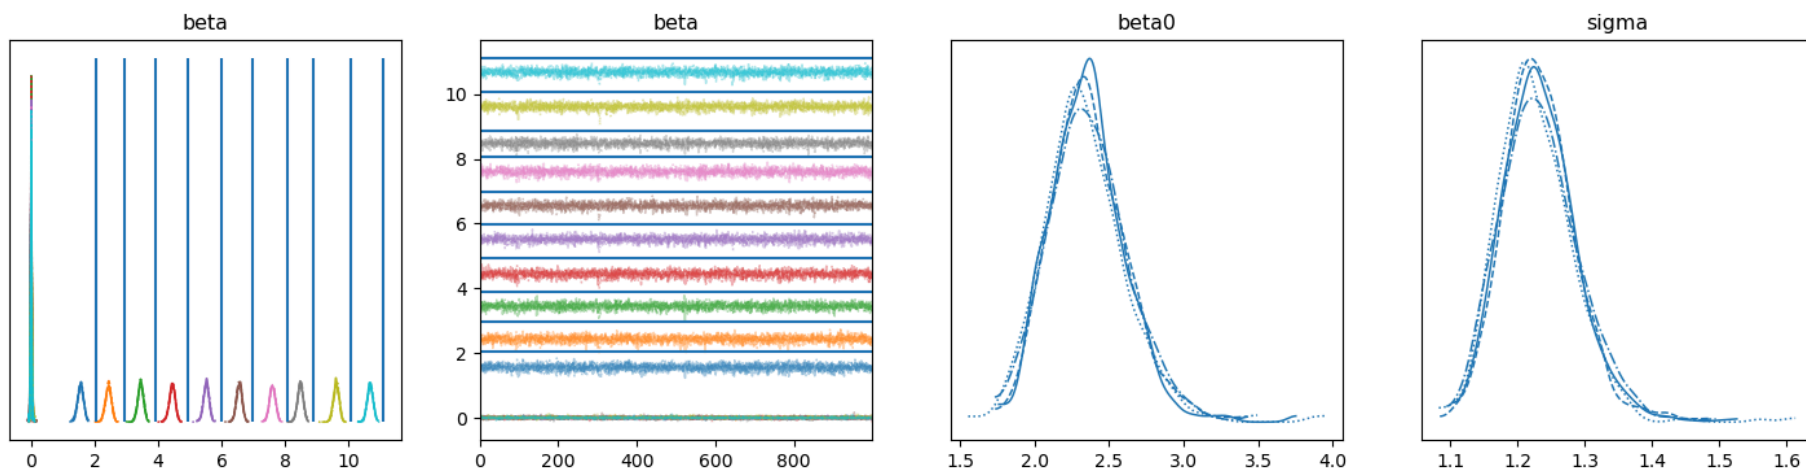
```
In [ ]: with pm.Model() as lasso:
    beta = pm.Laplace('beta', mu=0, b=1/75, shape=m, initval=betas[:,0])
    beta0 = pm.Normal('beta0', mu=0, sigma=10); sigma = pm.HalfNormal('sigma', sigma=100)
    y = pm.Normal('y', mu=beta0+pm.math.dot(X, beta), sigma=sigma, observed=y_obs)
    with lasso:
        idata4 = pm.sample()
```

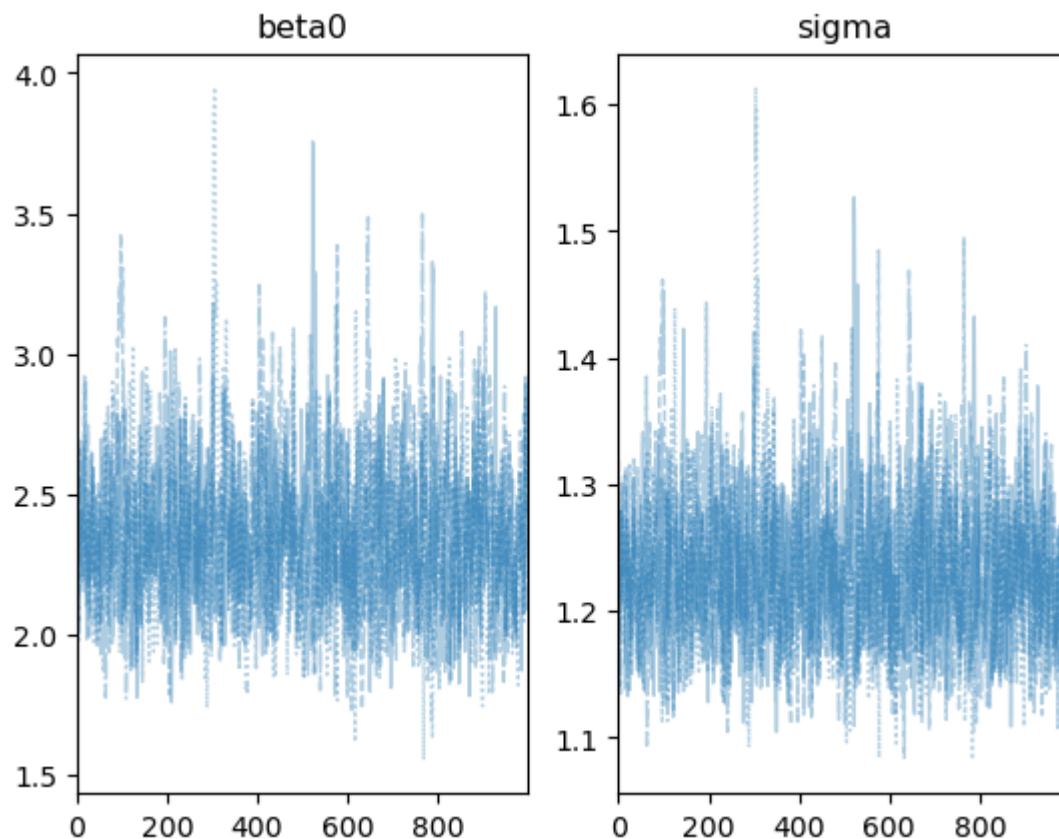
```
Auto-assigning NUTS sampler...  
Initializing NUTS using jitter+adapt_diag...  
Multiprocess sampling (4 chains in 4 jobs)  
NUTS: [beta, beta0, sigma]
```

```
100.00% [8000/8000 00:26<00:00 Sampling 4 chains, 0 divergences]
```

```
Sampling 4 chains for 1_000 tune and 1_000 draw iterations (4_000 + 4_000 draws total) took 26 seconds.
```

```
In [ ]: fig,ax = plt.subplots(1,4,figsize=(18,4)); fig2,ax2 = plt.subplots(1,2); az.plot_trace(idata4, axes=np.r_[ax[:3],ax
```





Bayesian Shrinkage Estimation: "The Horseshoe" [15 minutes]

The [PyMC overview](#) and [many other resources](#) provide **Horseshoe prior** [1] [2] implementations

| Half-Cauchy $HC_+(\xi)$ | Horseshoe Prior HSP | Shrinkage κ | Change of Variables | $|:-:|:-:|:-:|:-:|$ |

$$f(x \mid \xi) = \frac{2 \cdot 1_{[x \geq 0]}(x)}{\pi \xi \left[1 + \left(\frac{x}{\xi} \right)^2 \right]}$$

|

$$w_i | \tau \sim N(0, \sigma^2 = \lambda_i^2 \tau^2)$$

$$\lambda_i \sim HC_+(1)$$

$$\tau \sim HC_+(\tau_0)$$

|

$$\kappa_{\lambda_i} = 1/(1 + \lambda_i^2)$$

$$\lambda_i = \sqrt{1/\kappa_{\lambda_i} - 1}$$

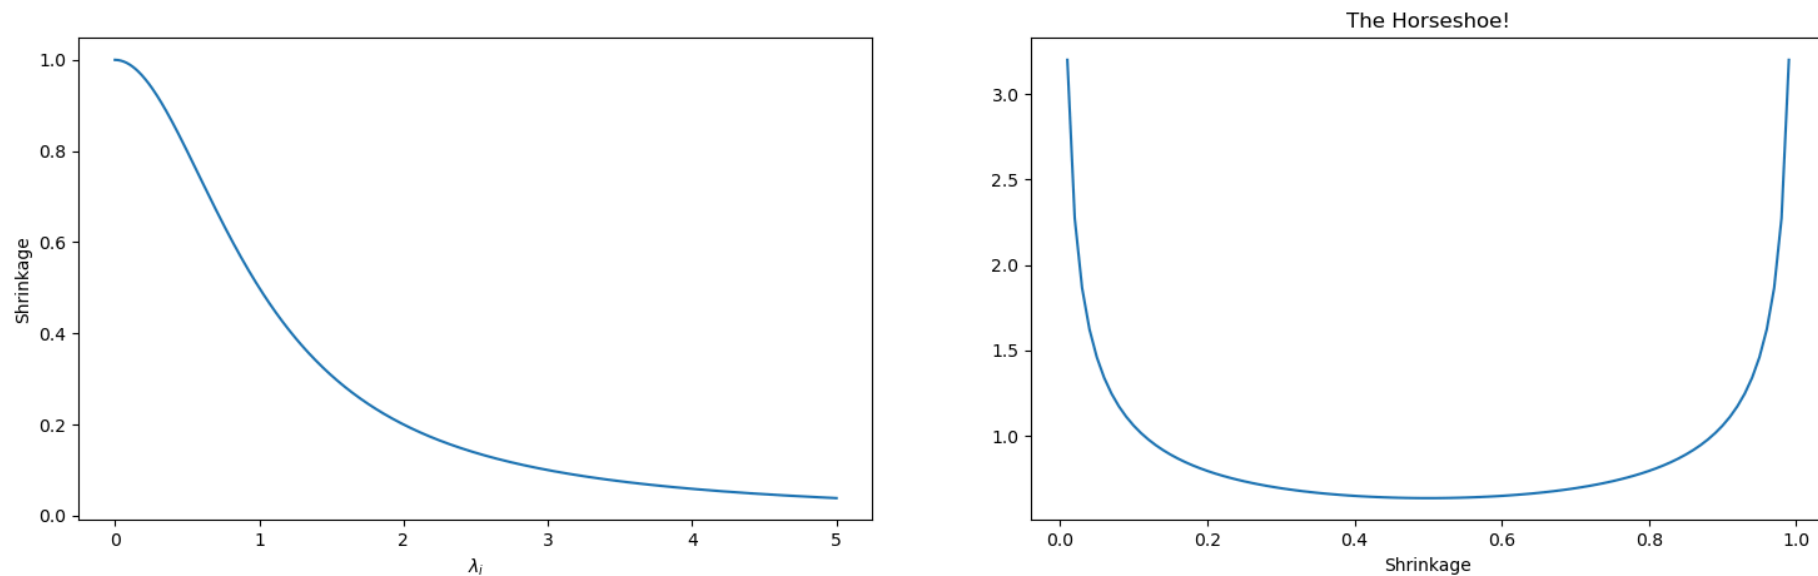
$$J_{\kappa_{\lambda_i}} = \frac{1}{2}(\kappa_{\lambda_i}^{-1} - 1)^{-\frac{1}{2}} \times \kappa_{\lambda_i}^{-2}$$

|

$$f(\kappa_{\lambda_i}) = f\left(\lambda_i = \sqrt{1/\kappa_{\lambda_i} - 1}\right) \times \underbrace{\frac{1}{2}(\kappa_{\lambda_i}^{-1} - 1)^{-\frac{1}{2}} \times \kappa_{\lambda_i}^{-2}}_{J_{\kappa_{\lambda_i}}}$$

|

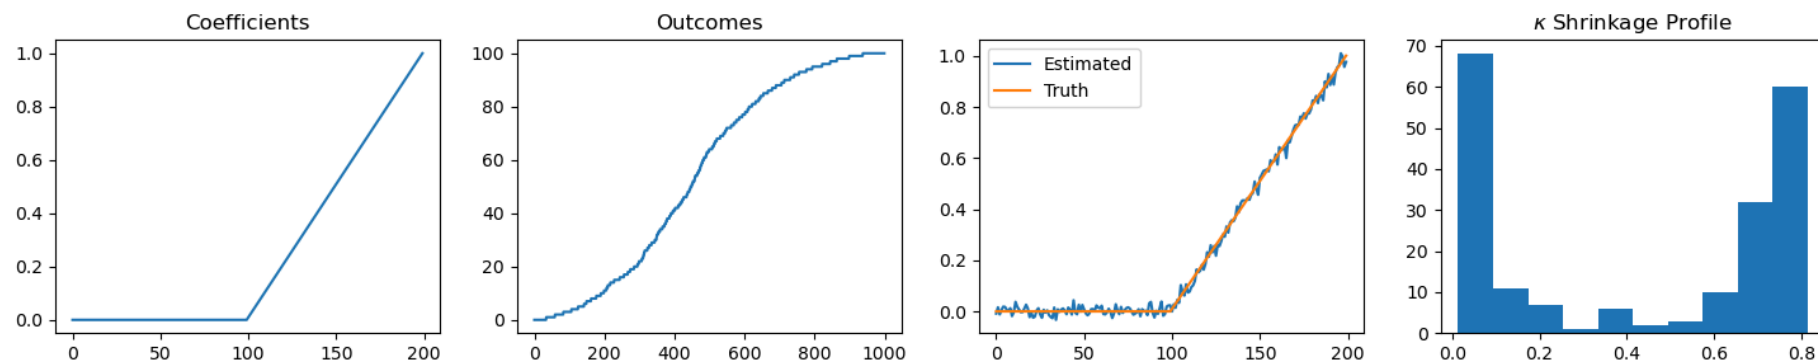
```
In [ ]: fig,ax = plt.subplots(1,2,figsize=(18,5)); support = np.linspace(0,5,1000); # shrnk = trans(sprpt) = 1/(1+sprt**2)
shrinkage = 1/(1+support**2); ax[0].plot(support, shrinkage); ax[0].set_ylabel("Shrinkage"); ax[0].set_xlabel("$\\lambda$")
# change of variables: sprpt = (1/shrnk-1)**0.5; E.g., 1/(1+.5**2), (1/.8-1)**0.5; jacobian: .5(1/shrnk-1)**(-.5)*s
shrinkage = np.linspace(0.01,.99,99); ax[1].plot(shrinkage, stats.halfcauchy(scale=1).pdf((1/shrinkage-1)**0.5) * .
```



Bayesian Shrinkage Estimation: "The Horseshoe" [15 minutes]

Trying it for **binomial regression** coefficients... a linear model and nonlinear transformation for non normal outcomes

```
In [ ]: from scipy.special import expit as invlogit; fig, ax = plt.subplots(1, 4, figsize=(18, 3)); K, Q = 100, 100; w = np.arange(1,
# nonzero and noise coefficients "weights" w (beta is reserved for the Half-Cauchy parameter below)
P = 1000; X = stats.bernoulli(p=0.5).rvs(size=(K+Q)*P).reshape(P, (K+Q)); N = 100; THETA = -25; x = stats.binom(p=invlogit(THETA
ax[2].plot(posterior.posterior['weights'].values.reshape((-1, 4000, 200))[0, :, :].mean(axis=0), label='Estimated'); ax
ax[3].hist(1/(1+posterior.posterior['lambdas'].values.reshape((-1, 4000, 200))[0, :, :].mean(axis=0)*2)); ax[3].set_ti
```



```
In [ ]: from pymc.math import invlogit as tt_invlogit
with pm.Model() as horseshoe:
    tau_0 = 1; tau = pm.HalfCauchy('tau', beta=tau_0, shape=1)
    lambdas = pm.HalfCauchy('lambdas', beta=1, shape=(Q+K)); theta = pm.Normal('theta', mu=0, sigma=50)
    weights = pm.Normal('weights', mu=0, sigma=tau*lambdas, shape=(Q+K))
    successes = pm.Binomial('successes', p=tt_invlogit(X@weights+THETA), n=[N]*P, observed=x)
    posterior = pm.sample()
```

Auto-assigning NUTS sampler...
 Initializing NUTS using jitter+adapt_diag...
 Multiprocess sampling (4 chains in 4 jobs)
 NUTS: [tau, lambdas, theta, weights]

100.00% [8000/8000 02:38<00:00 Sampling 4 chains, 831 divergences]

Sampling 4 chains for 1_000 tune and 1_000 draw iterations (4_000 + 4_000 draws total) took 159 seconds.
 The rhat statistic is larger than 1.01 for some parameters. This indicates problems during sampling. See <https://arxiv.org/abs/1903.08008> for details
 The effective sample size per chain is smaller than 100 for some parameters. A higher number is needed for reliable rhat and ess computation. See <https://arxiv.org/abs/1903.08008> for details
 There were 831 divergences after tuning. Increase `target_accept` or reparameterize.

Bayesian Shrinkage Estimation: "Regularized Horseshoe" [10 minutes]

$$\kappa_j = \frac{1}{1 + \lambda_j^2} \quad \text{for } \tau = \tau_0 = n = 1 \text{ and } \sigma^2 = \text{Var}(x_j) = s_{x_j}^2 \quad \text{generalizes too} \quad \frac{1}{1 + n\sigma^{-2}\tau^2 s_{x_j}^2 \lambda_j^2}$$

The *smaller* $\frac{n}{\sigma^2}\tau^2$ (with $s_{x_j}^2 = 1$) is the *greater* the number of the D parameters experiencing shrinkage [1] [2] [3]

- **Hyperparameter** τ_0 drives the **effective number of parameters** $E\left[\sum_{k=1}^D 1 - \kappa_j\right] = \frac{\frac{\tau_0}{\sigma}\sqrt{n}}{1 + \frac{\tau_0}{\sigma}\sqrt{n}} D = p_0$

- $\tau_0 = \frac{p_0}{D-p_0} \frac{\sigma}{\sqrt{n}}$ can be interpreted based on the number of "effectively non-zero" parameters p_0

```
In [ ]: import random; w_prime = w.copy(); w_prime[random.sample(range(len(w)), 150)] = 0 # zero out some coefficients
THETA_ = -6; x_ = stats.binom(p=invlogit(THETA_+X@w_prime), n=[N]*P).rvs(); #plt.figure(figsize=(18,3)); plt.plot(s
with pm.Model() as mod_reg_1p0:
    tau_0 = 1; tau = tau_0 #tau = pm.HalfCauchy('tau', beta=tau_0, shape=1)
    lambdas = pm.HalfCauchy('lambdas', beta=1, shape=(Q+K)); theta = pm.Normal('theta', mu=0, sigma=20)
    weights = pm.Normal('weights', mu=0, sigma=tau*lambdas, shape=(Q+K))
    successes = pm.Binomial('successes', p=tt_invlogit(X@weights+theta), n=[N]*P, observed=x_)
    posterior_reg_1p0 = pm.sample()
with pm.Model() as mod_reg_0p03:
    tau_0 = 0.03; tau = tau_0 #tau = pm.HalfCauchy('tau', beta=tau_0, shape=1)
    lambdas = pm.HalfCauchy('lambdas', beta=1, shape=(Q+K)); theta = pm.Normal('theta', mu=0, sigma=20)
    weights = pm.Normal('weights', mu=0, sigma=tau*lambdas, shape=(Q+K))
    successes = pm.Binomial('successes', p=tt_invlogit(X@weights+theta), n=[N]*P, observed=x_)
    posterior_reg_0p03 = pm.sample()
```

Auto-assigning NUTS sampler...

Initializing NUTS using jitter+adapt_diag...

Multiprocess sampling (4 chains in 4 jobs)

NUTS: [lambdas, theta, weights]

100.00% [8000/8000 03:32<00:00 Sampling 4 chains, 1,158 divergences]

Sampling 4 chains for 1_000 tune and 1_000 draw iterations (4_000 + 4_000 draws total) took 213 seconds.

The rhat statistic is larger than 1.01 for some parameters. This indicates problems during sampling. See <https://arxiv.org/abs/1903.08008> for details

The effective sample size per chain is smaller than 100 for some parameters. A higher number is needed for reliable rhat and ess computation. See <https://arxiv.org/abs/1903.08008> for details

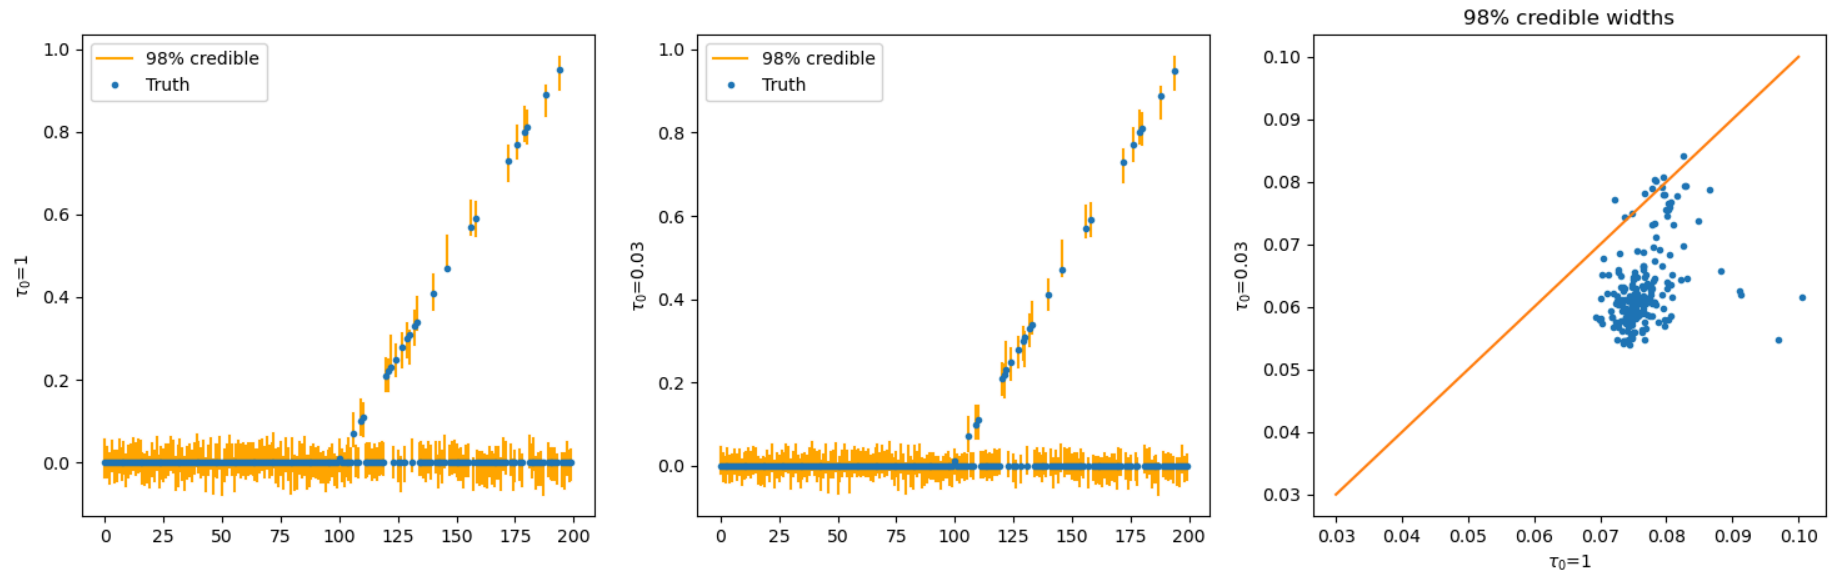
There were 1158 divergences after tuning. Increase `target_accept` or reparameterize.

Bayesian Shrinkage Estimation: "Regularized Horseshoe" [5 minutes]

```
In [ ]: # Approximating the calculations of the previous page which are technically for linear model regression
sigma2hat = \
(x_ - N*invlogit(posterior_reg_1p0.posterior['theta'].values.reshape((-1,4000,1))[0].mean(axis=0)+
                X@posterior_reg_1p0.posterior['weights'].values.reshape((-1,4000,200))[0].mean(axis=0))).var()
tau0 = ((w_prime!=0).sum()/(Q+K-(w_prime!=0).sum()))*(sigma2hat/(X.var(axis=0).mean()*1000))*0.5; tau0
#tau0*(X.var(axis=0).mean()*1000)/sigma2hat)**0.5/(1+tau0*(X.var(axis=0).mean()*1000)/sigma2hat)**0.5)*200
```

Out []: 0.027886390597189

```
In [ ]: fix,ax = plt.subplots(1,3,figsize=(18,5)); par_1p0 = np.array(2*[np.arange(len(posterior_reg_1p0.posterior['weights
```



Homework 6: Part III Robust regression: scale mixtures (of normals)

$$\int \frac{w\lambda_i}{\sqrt{2\pi}} e^{-\frac{1}{2}(w^{-2}\lambda_i^{-2}(y_i-\mu)^2)} \frac{\frac{\nu}{2}}{\Gamma(\frac{\nu}{2})} \lambda_i^{\nu-\frac{1}{2}} e^{-\frac{\nu}{2}\lambda_i^{-2}} d\lambda_i = \frac{\Gamma\left(\frac{\nu+1}{2}\right)}{\Gamma(\frac{\nu}{2})\sqrt{\pi\nu}w^2} \left(1 + \frac{1}{\nu}\left(\frac{y_i-\mu}{w}\right)^2\right)^{-\frac{\nu+1}{2}}$$

$$\begin{aligned} y_i|\lambda_i &\sim \mathcal{N}(X\beta, \sigma^2 = w^2\lambda_i^2) \\ \lambda_i^{-2} &\sim \text{Gamma}(\alpha = \nu/2, \beta = \nu/2) \implies y_i \sim t_\nu(\mu, w^2) \end{aligned}$$

1. Return to your kaggle.com regression data set; or, find another data set; and use the above specification to perform a robust regression analysis in `PyMC`
2. Use the posterior distributions of the λ_i 's to identify "outlier" (and potentially "influential") data points
3. [Optional] Assess the performance of the MCMC and any issues or warnings in the [standard manner](#)
4. [Optional] Perform **Multiple Linear Regression** diagnostics... residual plots, etc.